

love this?

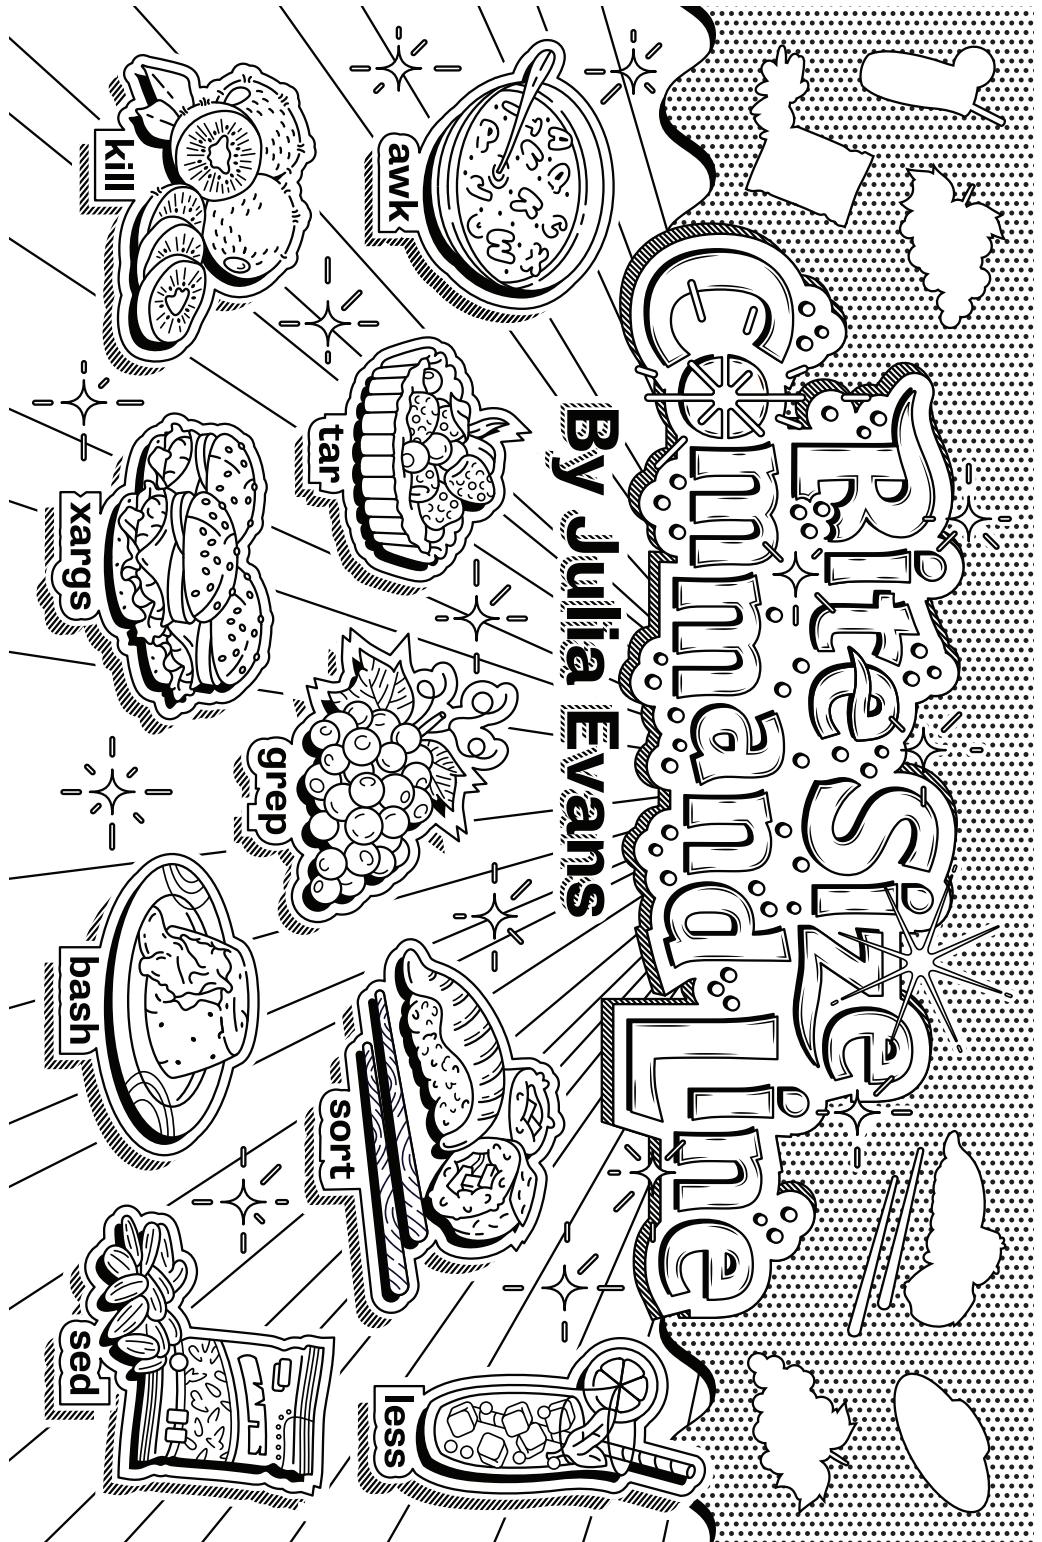
find more awesome zines at

→ [jvns.ca/zines](http://jvns.ca/zines) ←

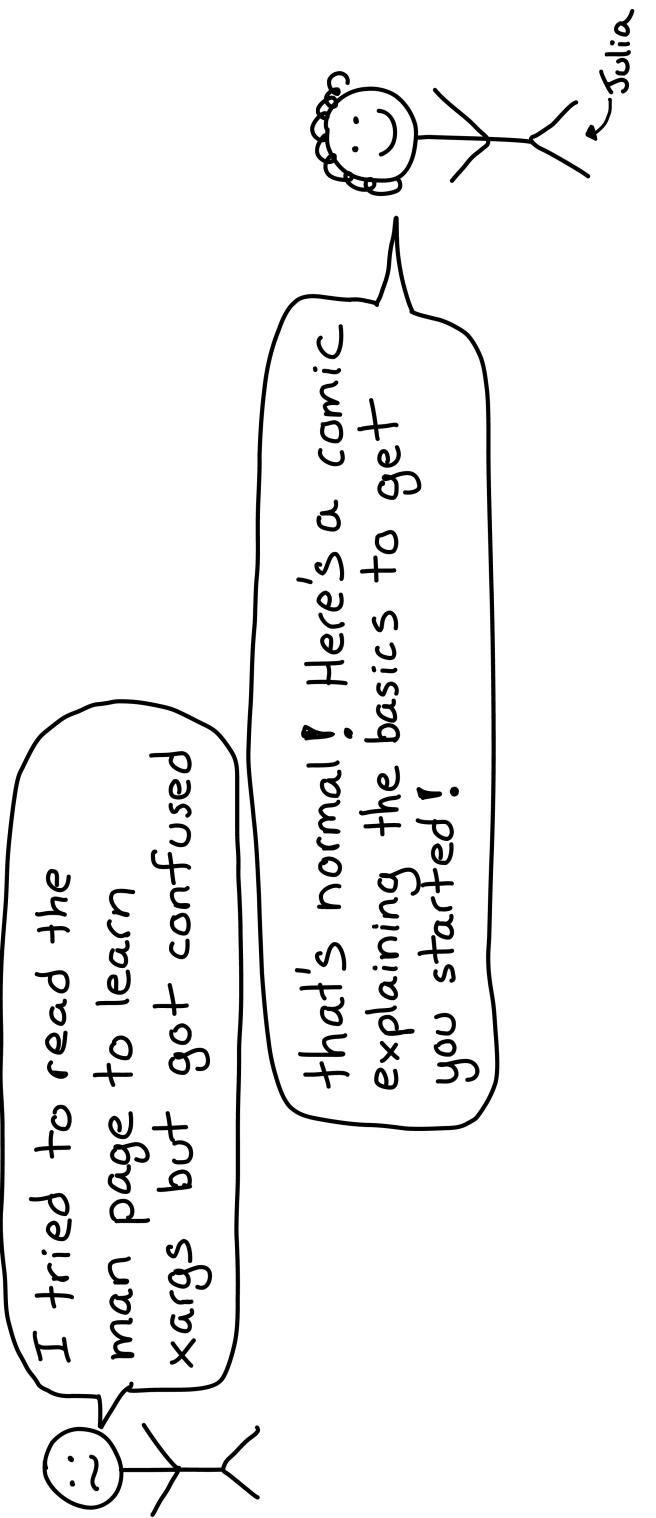
CC-BY-NC-SA computer wizard industries 2018

# Rite Size Mind Line

By Julia Evans



This zine explains some of the most useful Unix command line tools in 1 page each.



Even if you've used the command before, I might have a new trick or two for you ❤

\*

## more useful tools

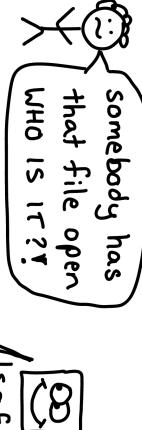
- make
- jq
- nohup
- disown
- cut/paste
- sponge
- xxd
- hexdump
- objdump
- strings
- screen
- tmux
- date
- entr
- seq
- join
- parallel:
  - GNU parallel
  - pigz / pixz
  - sort --parallel

\*

# lsof

`lsof`

stands for list open files



what `lsof` tells you  
for each open file:

- pid
- file type (regular? directory?  
FIFO? socket?)
- file descriptor (FD column)
- user
- filename / socket address

`-P PID`  
list the files PID has  
open

`lsof /some /dir`  
list just the open files  
in `/some /dir`

## find deleted files

`$ lsof | grep deleted`

will show you deleted files!

examples:  
`-i -n -P` → `-n & -P` mean  
`-i :8080` "don't resolve  
`-i TCP` host names / ports"  
`-i -s TCP :LISTEN`

You can recover open deleted  
files from  
`/proc/<pid>/fd/<fd>`

process that opened the file

## netstat

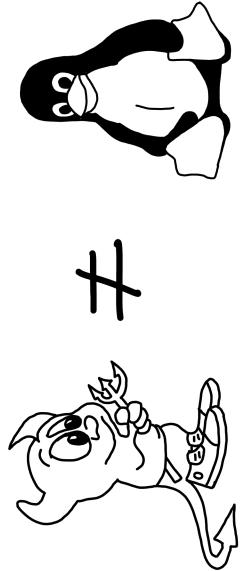
another way to list open  
sockets on Linux is:

`netstat -tunapl`

↑  
tuna, please!  
On Mac netstat has  
different args.

# Table of contents

BSD≠GNU.....	4	bash tricks.....	10-11	misc commands...17
grep.....	5	disk usage.....	12	head & tail..... 18
find.....	6	tar.....	13	less..... 19
xargs.....	7	ps.....	14	kill .....
awk.....	8	top.....	15	cat..... 21
sed.....	9	sort & uniq.....	16	lsof..... 22

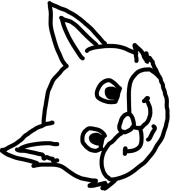


For almost all these tools, there are at least 2 versions:

① The BSD version (on BSDs & Mac OS)

② The GNU version (on Linux)

All of the examples in this zine were tested on Linux. Some things like sed -i are different on Mac. Be careful when writing cross-platform scripts! You can install the GNU versions on Mac with `brew install coreutils`.



## cat & friends

21

you can use cat as an EXTREMELY BASIC text editor:

```
$ cat myfile.txt
prints contents of myfile.txt
$ cat *.txt
prints all .txt files put together!
```

`cat -n`  
prints out the file with line numbers!

- 1 Once upon a midnight...
- 2 Over many a quaint...
- 3 While I nodded, nearly

## tee

'tee file.txt' will write its stdin to both stdout and file.txt  
 $\xrightarrow{\text{stdin}}$  tee a.txt  
 $\xrightarrow{\text{stdout}}$  a.txt

## zcat

cats a gzipped file!  
 Actually just a 1-line shell script that runs `gzip -cd`, but easier to remember.

# Kill

kill doesn't just kill programs

 You can send ANY signal to a program with kill!

**kill - SIGNAL PID**

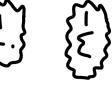
name or number

**killall - SIGNAL NAME**

signals all processes called NAME for example

\$ killall firefox

useful flags:

 wait for all signaled processes to die

 ask before signalling

<u>name</u>	<u>num</u>
kill	SIGTERM 15
kill -q	SIGKILL 9
kill -KILL	10 USR1 11 SEGV 12 USR2
kill -HUP	SIGHUP 1
kill -STOP	SIGSTOP 19
kill -TSTP	SIGTSTP 18
kill -CONT	SIGCONT 19
kill -TRIN	SIGTRAP 5
kill -TTOU	SIGTTOU 21
kill -URG	SIGURG 23
kill -XCPU	SIGXCPU 24
kill -WINCH	SIGWINCH 28
kill -POLL	SIGPOLL 29
kill -PWR	SIGPWR 30
kill -SYS	SIGSYS 31

Kill -R	
lists all signals	
1 HUP	2 INT
2 INT	3 QUIT
3 QUIT	4 ILL
4 ILL	5 TRAP
5 TRAP	6 ABRT
6 ABRT	7 BUS
7 BUS	8 FPE
8 FPE	9 KILL
9 KILL	10 USR1
10 USR1	11 SEGV
11 SEGV	12 USR2
12 USR2	13 PIPE
13 PIPE	14 ALRM
14 ALRM	15 TERM
15 TERM	16 STKFLT
16 STKFLT	17 CHLD
17 CHLD	18 CONT
18 CONT	19 STOP
19 STOP	20 TSTP
20 TSTP	21 TTIN
21 TTIN	22 TTOU
22 TTOU	23 URG
23 URG	24 XCPU
24 XCPU	25 XFSZ
25 XFSZ	26 VTIME
26 VTIME	27 PROF
27 PROF	28 WINCH
28 WINCH	29 POLL
29 POLL	30 PWR
30 PWR	31 SYS

# grep

5

grep lets you search files for text

\$ grep bananas foo.txt

Here are some of my favourite grep command line arguments!

**-i** case insensitive

**-A** show context for your search.

**-B** will show 3 lines of context after a match

**-C** aka egrep

**-E** aka egrep

**-F** don't treat the string as a regex

string as a regex

**-l** only show the filenames of the files that matched

**-r** recursive! search all the files in a directory.

**-o** only print the matching part of the line (not the whole line)

**-a** search binaries: treat binary data like it's text instead of ignoring it!

**-g** grep alternatives

**-g** ripgrep (better for searching code!)

# find

6

**find** searches a directory for files  
 find /tmp -type d, -print  
 ↗  
 directory which files action to do with the files  
 to search

**-name/-iname**  
 case insensitive  
 the filename! eg  
 -name '\*.txt'

**-path / -i-path**  
 search the full path!  
 -path '/home/\*/\*.go'

**-size 0**

find empty files!  
 Useful to find files you created by accident

**-exec COMMAND**

action: run COMMAND on every file found

**-type [TYPE]**  
 f: regular file l: symlink  
 d: directory + more!

**-maxdepth NUM**

only descend NUM levels when searching a directory

**locate**

The locate command searches a database of every file on your system.  
 good: faster than find  
 bad: can get out of date  
 \$ sudo updatedb  
 updates the database

# less

19

**less is a pager**  
 that means it lets you view (not edit) text files or piped in text  
 man uses your pager (usually less) to display man pages

many vim shortcuts work in less  
 / search  
 n/N next / prev match  
 j/K down/up a line  
 m/, mark/return to line  
 g/G beginning/end of file  
 {gg in vim}

**q**  
 quit ↴

**V ~ lowercase**  
 edit file in your \$EDITOR

arrow keys, Home/End,  
 Pg Up, Pg Dn work in less

**less -r**  
 displays bash escape codes as colours  
 try ls --color | less -r  
without -r  
 a.txt ↴ a.txt  
 ESC C 0 m ESC C 0 i 3 l m a . t x t g 2 red, bold ESC C 0 m

**F**  
 press F to keep reading from the file as it's updated (like tail -f)

**+ runs a command when less starts**  
 less +F : follow updates  
 less +G : start at end of file  
 less +20% : start 20% into file  
 less +/foo : search for 'foo' right away

# head & tail

**head**  
shows you the first 10 lines of a file

if you pipe a program's output to head, the program will stop after printing 10 lines (it gets sent `SIGPIPE`)

## -c NUM

show the first/last NUM bytes of the file

head -c 1k

will show the first 1024 bytes

## tail

**tail** shows the last 10 lines!

`tail -f FILE` will follow:  
print any new lines added to the end of FILE. Super useful for log files!

## tail --retry

keep trying to open file if its inaccessible

## tail --pid PID

stop when process PID stops running (with -f)

## -n NUM

**-n NUM** (either head or tail) will change the # lines shown

`head -n -NUM` } show all tail -n +NUM } but the last / first NUM lines

## tail --follow=name

Usually tail -f will follow a file descriptor.

`tail --follow=name FILENAME` will keep following the same filename, eg if the file descriptor is rotated.

# Xargs

7

this is useful when you want to run the same command on a list of files!

- delete (`xargs rm`)
- combine (`xargs cat`)
- search (`xargs grep`)
- replace (`xargs sed`)

how to replace "foo" with "bar" in all .txt files:

```
find . -name '*.txt' |  
xargs sed -i $'foo/bar/g
```

how to lint every Python file in your Git repo:

```
git ls-files | grep .py |  
xargs pep8
```

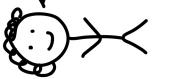
more useful xargs options

`(-n 1)` makes xargs run max-args for each input.  
`(-P)` is the max number of parallel processes xargs will start

# awk

8

awk is a tiny programming language for manipulating columns of data



I only know how to do 2 things with awk but it's still useful!

basic awk program structure

```
BEGIN { ... }
CONDITION {action}
CONDITION {action}
END {...}
    do action on lines matching condition
```

so MANY unix commands print columns of text (ps! ls!)

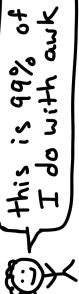
so being able to get the column you want with awk is GREAT

awk program example: sum the numbers in the 3rd column

```
action
{ s += $3 }
END { print s }
    at the end, print the sum!
```

extract a column of text with awk

```
awk -F, ' {print $5}'
```



single quotes! 5<sup>th</sup> column separator  
this is 99% of what I do with awk

awk program example: print every line over 80 characters

```
length($0) > 80
    condition
(there's an implicit {print} as the action)
```

# misc commands



17

**pv**

"pipe viewer", gives you stats on data going through a pipe

**cail**

a tiny calendar ☰

**ncdu**

figure out what's using all your disk space

**ts**

add a timestamp in front of every input line

**comm**

find lines 2 sorted files have in common

**rlwrap**

adds history & ctrl support to REPLs that don't already have them  
(rl stands for readline)

```
$ rlwrap python
```

**xsel / xclip**

copy/paste from system clipboard.  
(pbcopy / pbpaste on Mac)

copy/paste from system clipboard.  
(pbcopy / pbpaste on Mac)

# sort & uniq

sort sorts its input  
\$ sort names.txt  
the default sort is alphabetical.

uniq removes duplicates  
a  
b  
b => a's ! uniq  
a  
c  
c  
notice there are still 2  
only uniqueness adjacent matching lines

sort -n numeric sort  
'sort' order 'sort-n' order  
12 12  
15000 48  
~ 48  
~ 6020 96  
~ 6020  
96 15000

sort + uniq = ♥  
Pipe something to 'sort | uniq' and you'll get a deduplicated list of lines ! sort -u does the same thing.  
b | sort -u => a  
b | sort -u => b  
a or sort | uniq

sort -h: human sort  
'sort-n' order 'sort-h' order  
15G 45K  
30M 30M  
~ 45K 15G  
~ 200G 200G

useful example:  
du -sh \* | sort -h

# Sed

q

some more sed incantations...

sed is most often used for replacing text in a file  
\$ sed s/cat/dog/g file.txt  
can be a regular expression

change a file in place with -i  
in GNU sed it's -i  
in BSD sed, -i SUFFIX  
confuses me every time.

sed s+cat/+dog/+  
use + as a regex delimiter  
way easier than escaping / is like s/cat\/\|/dog\|/ !

sed -n 12 p  
print 12<sup>th</sup> line

sed 6  
double space a file (good for long error lines)

sed 'cat/a dog'  
append 'dog' after lines containing 'cat'

sed -n 5,30 p  
print lines 5-30

# bash tricks

10

## \* **ctrl + r \***

search your history!

I use this ❤ constantly ❤ to rerun commands

\$ convert file.{jpg,png}  
expands to  
  
\$ convert file.jpg file.png  
 $\{1..5\}$  expands to 1 2 3 4 5  
(for i in {1..100};)

## \* magical braces \*

\$ convert file.{jpg,png}

expands to

\$ convert file.jpg file.png  
 $\{1..5\}$  expands to 1 2 3 4 5  
(for i in {1..100};)

## loops

```
for i in *.png
```

```
do
```

```
convert $i $i.jpg
```

```
done
```



gives the output of a command

```
$ touch file-$(date -I)
```

create a file named file-2018-05-25

!!

expands to the last command run  
\$ sudo !!

commands that start with a space don't go in your history. good if there's a password

## more keyboard shortcuts

ctrl a beginning of line

ctrl e end of line

ctrl l clear the screen

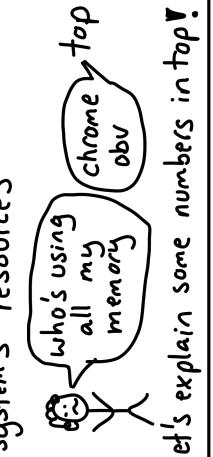
+ lots more emacs  
shortcuts too!

15

# top

## load average

3 numbers that roughly reflect demand for your CPUs on the system in the last 1, 5, and 15 minutes if it's higher than the # of CPUs you have, that's often bad



## top

a live-updating summary of the top users of your system's resources



## % CPU



this column is given as % of a single core. If you have 4 cores, this can go up to 400 %!

## memory

4 numbers:

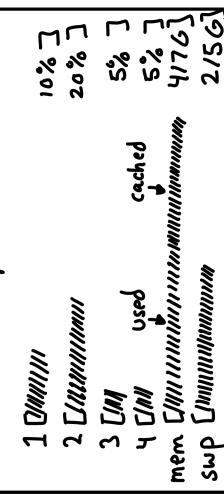
total / free / used / cached

One perhaps unexpected thing:  
total is not free + used!  
total = free + used + cached  
filesystem cache

15

## htop

a prettier & more interactive version of top \*



## RES

this column is the "resident set size" aka how much RAM your process is using.

SHR is how much of the RES is shared with other processes

# PS

## ps

ps shows which processes are running  
I usually run ps like this:

```
$ ps aux
```

u means include  
username column  
(ps -ef works too)

a+\* together  
show all process

## w

is for wide. ps auxww will show all the command line args for each process

## e

is for environment. ps auxe will show the environment vars!

## \* process state \*

Here's what the letters in ps's STATE column mean:

R: running  
S/D: asleep  
Z: zombie  
I: multithreaded  
+: in the foreground

## f

is for "forest" !. ps auxf will show you an ASCII art process tree!

ps tree can display a process tree too

you can choose which columns to show with ps (ps -eo ...) One cool column is 'wchan' which tells you the name of the kernel function if the process is sleeping  
try it:  
ps -eo user,pid,wchan,cmd

## wchan

ps has 3 different sets of command line arguments ↪  
1. UNIX ( 1 dash)  
2. BSD (no dash)  
3. GNU ( 2 dashes)

You can write monstrities like:  
\$ ps f -f ↪ full format  
forest(BSD) (UNIX)

# more bash tricks 11

## cd -

changes to the directory you were last in

pushd & popd let you keep a stack

## ctrl+z

suspends (SIGTSTP) the running program

## fg

brings backgrounded/suspended program to the foreground

## bg

starts suspended program & backgrounds it (use after ctrl+z)

## shellcheck

shell script linter! helps spot common mistakes.

## <()

process substitution  
treat process output like a file (no more temp files!)  
eg:  
\$ diff <(ls) <(ls -a)

## fc

"fix command"  
open the last command you ran in an editor  
then run the edited version

## type

tells you if something is a builtin, program, or alias  
try running type on ~~{time}~~ ~~ping~~ ~~[pushd]~~  
(they're all different types)

# disk usage

12

## du

tells you how much disk space files / directories take up  
-S summary: total size of all files in a directory  
-h human readable sizes

## df -i

instead of % disk free, report how many inodes are used / free on each partition  
running out of inodes is VERY ANNOYING - You can't create new files!

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/sda3	18G	6G	2.5G	86%	/
udev	4.83M	4.0K	4.83M	1%	/dev
tmpfs	99M	1.4M	97M	2%	/run
/dev/sda4	167G	157G	9.9G	95%	/home

## ncdu

see what's using disk space navigate with arrow keys  
17.5 GiB [#####] /music  
3.2 GiB [##] /pictures  
5.7 MiB [ ] /text  
2.0 MiB [ ] file.pdf

iostat	
get statistics about disk reads / writes	interval to report at
# iostat 5	
Device: kB_read/s kB_wrtn/s	
sda 2190.21 652.87	
sdb 6.00 0.00	

# tar

13

The tar file format combines many files into one file.  
a.txt dir/c.txt  
b.txt  
.tar files aren't compressed by themselves. Usually you gzip them: .tar.gz or .tgz!

Usually when you use the 'tar' command, you'll run some incantation To unpack a tar.gz, use:  
-tar -xzf file.tar.gz  
let's learn!  
what's xzf?

-X is for extract into the current directory by default (change with -C)

-C is for create makes a new tar file!

-t is for list lists the contents of a tar archive

-f is for file which tar file to create or unpack & more! see the man page :)

tar can compress / decompress

-Z gzip format (.gz)  
-j bzip2 format (.bz2)  
-J xz format (.xz)  
\$ tar -tzvf file.tar.bz2  
\$ tar -tvf file.tar.gz  
\$ tar -czf file.tar.gz  
files to go in the archive