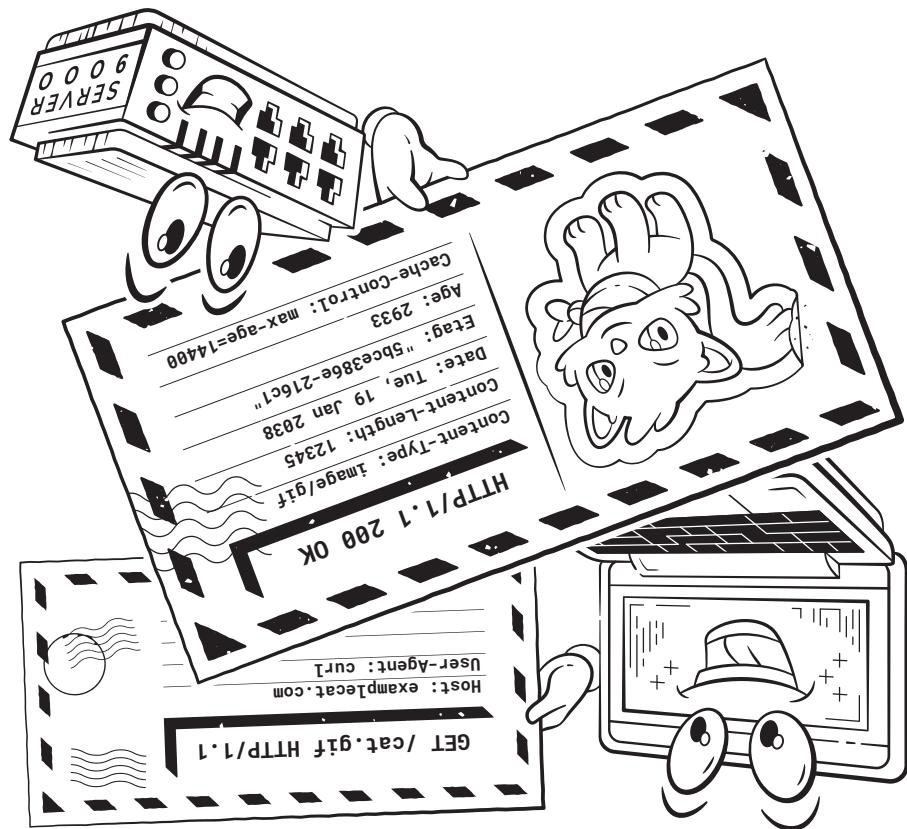


by Julia Evans

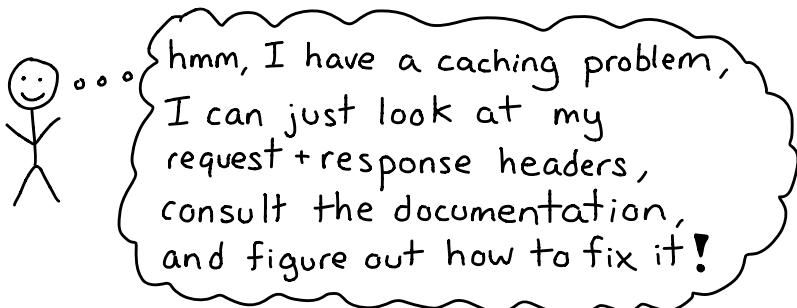
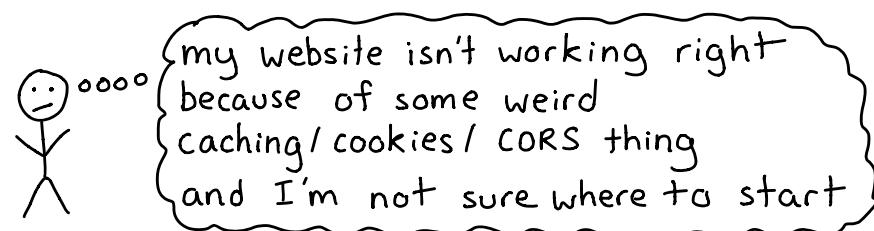


Like this?  
more Zines at  
[wizardzines.com](http://wizardzines.com)

# about this zine

Your browser uses HTTP every time it visits a website. Like a lot of the tech that runs the internet, understanding HTTP isn't that hard!

This zine's goal is to take you from:



## credits

Cover art: Vladimir Kašiković  
Editing: Dolly Lanuza, Kamal Marhubi  
special thanks to Marco Rogers for suggesting the idea of a HTTP zine

# how to learn more

## ♥ Mozilla Developer Network

<https://developer.mozilla.org>

MDN is a fantastic wiki maintained by Mozilla. It has tutorials and reference documentation for HTML, CSS, HTTP, Javascript. It's the best place to start for reference documentation.

## ♥ OWASP

<https://cheatsheetseries.owasp.org>

OWASP is an organization that publishes security best practices. If you have a question about web security, they've probably published a cheat sheet or guide to help you.

## ♥ httpstatuses.com

Nice little site that explains all the HTTP status codes.

## ♥ RFCs

<https://tools.ietf.org/html/rfcXXXX>

put RFC number here

RFCs are numbered documents (like "RFC 2631"). Every Internet protocol (like TLS or HTTP) has an RFC. These are where you go to find the Official Final Answers to technical questions you have about any internet standard. The HTTP standard is mostly documented in 6 RFCs numbered 7230 to 7235.



is the Host header actually required?



Yes, section 5.4 of RFC 7230 says so!

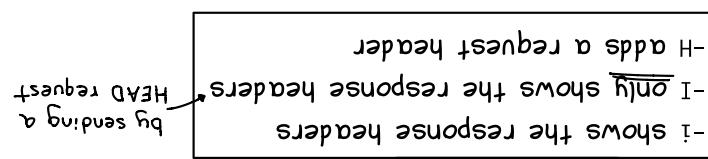
the final answer  
Don't be scared of using an RFC if you want to know for sure!

4	What's HTTP?
5	How URLs work
6	What's a header?
7	Anatomy of an HTTP request
8-9	Request methods (GET! POST!)
10	Request headers
11	Using HTTP APIs
12	Anatomy of an HTTP response
13	Response headers
14	Status codes (200! 404!)
15	How cookies work
16-17	Content delivery networks & caching
18	Redirections
19	HTTP/2
20-21	HTTPS & certificates
22-24	Same origin policy & CORS
25	Security headers
26-27	Exercises & how to learn more

## Table of Contents

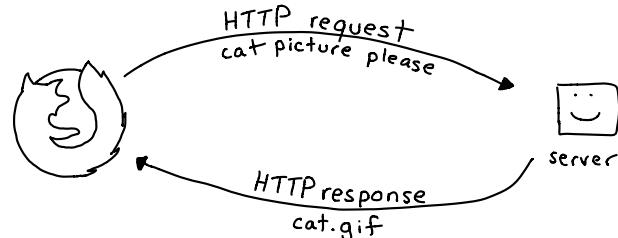
## HTTP exercises

Making HTTP requests with curl to real internet websites and trying different headers is my favourite way to play around with HTTP & learn.

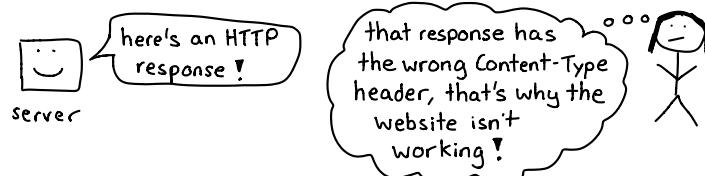


# what's HTTP?

HTTP is the protocol (Hypertext Transfer Protocol) that's used when you visit any website in your browser.



The exciting thing about HTTP is that even though it's used for literally every website, HTTP requests and responses are easy to look at and understand:



Example of what an HTTP request and response might look like:

request	response
request line headers	{ GET / HTTP/1.1 status
{ Host: examplecat.com User-Agent: curl Accept: */*	{ HTTP/1.1 200 OK headers
	{ Cache-Control: max-age=604800 Content-Type: text/html Etag: "1541025663+ident" Server: ECS (nyb/1D0B) Vary: Accept-Encoding X-Cache: HIT Content-Length: 1270
body	{ <!doctype html> <title>Example Cat</title> ...

All that text is a lot to understand, so let's get started learning what all of it means!

# security headers

These are headers your server can set. They ask the browser to protect your users' data against attackers in different ways:

**Content-Security-Policy** often called CSP

Only allow CSS / Javascript from certain domains you choose to run on your website. Helps protect against cross-site-scripting (aka XSS) attacks.

**Referrer-Policy**

Control how much information is sent to other sites in the Referrer header. Example: Referrer-Policy: no-referrer.  
spelling is inconsistent with Referer header !!

**Strict-Transport-Security** often called HSTS

Require HTTPS. If you set this the client (browser) will never request a plain HTTP version of your site again. Be careful!  
You can't take it back!

**Expect-CT**

Certificate Transparency (CT) is a system that can help find malicious SSL certificates issued for your site. This header gives the browser a URL to use to report bad certificates to you.

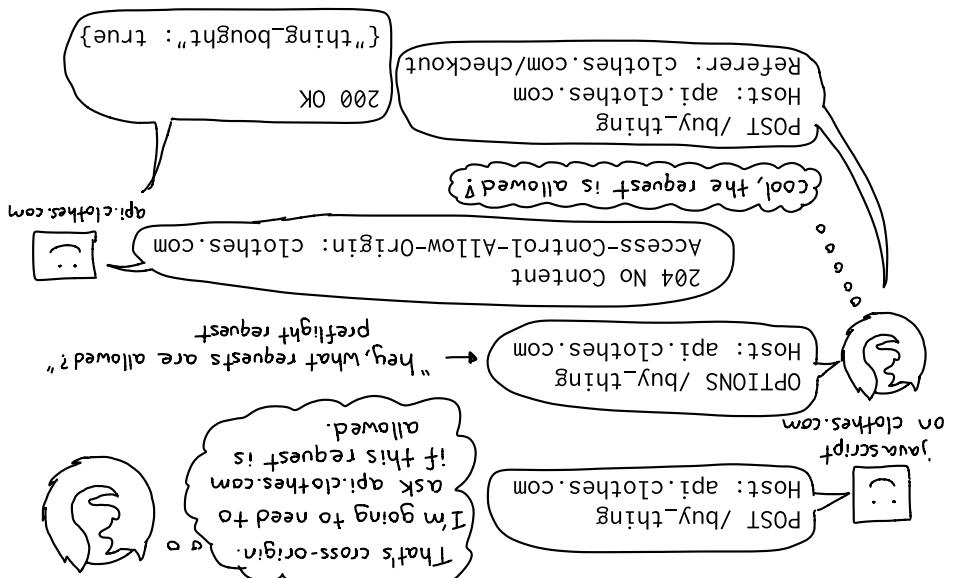
**X-XSS-Protection**

Another way to protect against XSS attacks. Not supported by all browsers, Content-Security-Policy is more powerful.

fragment id	This isn't sent to the server at all. It's used either to jump to an HTML tag ( <code>&lt;a id="banana".&gt;</code> ) or by #banana
space is %20, % is %25, etc.	
% + hex representation of ASCII value.	
URL you need to percent encode them as characters like spaces, @, etc. So to put them in a URL you need to have certain special characters like spaces, @, etc.	
URLs aren't allowed to have certain special characters like spaces, @, etc.	
name = value separated by &	
hair=short&color=black&name=m%20darcy	
gray cat@). Example:	
a different version of a page ("I want a light gray cat"). Example:	
Query parameters are usually used to ask for a different version of a page ("I want a light gray cat"). Example:	
like: GET /cats?color=light%20gray HTTP/1.1	
query parameters are combined in the request, path to ask the server for. The path and the path	
Defaults to 80 for HTTP and 443 for HTTPS.	
domain	Where to send the request. For HTTPS(s), examplecat.com the Host header gets set to this (Host: example.cat)
scheme	Protocol to use for the request. Encrypted (https), insecure (http), or something else entirely (ftp).
port	Default to 80 for HTTPS and 443 for HTTP.
path	/cats
query	Path to ask the server for. The path and the path
parameters	Query parameters are combined in the request, path to ask the server for. The path and the path
URL	name = value separated by &
encoding	hair=short&color=black&name=m%20darcy
space is %20	gray cat@). Example:

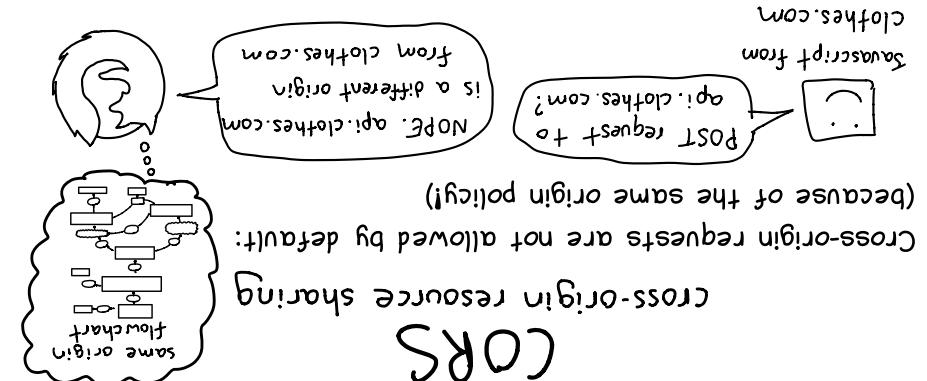
## HOW URLs WORK

This OPTIONS request is called a "preflight" request, and it only happens for some requests, like we described in the diagram on the same-origin policy page. Most GET requests will just be sent by the browser without a preflight request first, but POST requests that send JSON need a preflight.



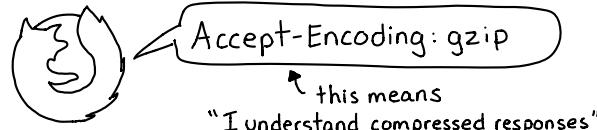
Here's what happens:

If you run `api.clothes.com`, you can allow `clothes.com` to make requests to it using the `Access-Control-Allow-Origin` header.

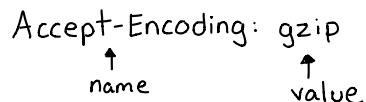


# what's a header?

Every HTTP request and response has headers. Headers are a way for the browser or server to send extra information!



Headers have a name and a value.



Header names aren't case sensitive:

totally valid → aCcEPT-eNcOdInG : gzip

There are a few different kinds of headers:

Describe the body:

Content-Type: image/png    Content-Encoding: gzip  
Content-Length: 12345    Content-Language: es-ES

Ask for a specific kind of response:

Accept: image/png    Accept-Encoding: gzip  
Range: bytes=1-10    Accept-Language: es-ES

Every Accept-  
header has a  
corresponding  
Content- header

Manage caches:

ETag: "abc123"  
If-None-Match: "abc123"  
Vary: Accept-Encoding  
If-Modified-Since: 3 Aug 2019 13:00:00 GMT  
Last-Modified: 3 Feb 2018 11:00:00 GMT  
Expires: 27 Sep 2019 13:07:49 GMT  
Cache-Control: public, max-age=300

Say where the request comes from:

User-Agent: curl    Referer: https://examplecat.com

Cookies:

Set-Cookie: name=julia; HttpOnly (server → client)  
Cookie: name=julia (client → server)

6 and more!

# why the same origin policy matters

Browsers work hard to make sure that evil.com can't make requests to other-website.com. But evil.com can request other-website.com from its own server, what's the big deal?

2 reasons it's important to restrict Javascript on websites from making arbitrary requests from your browser:

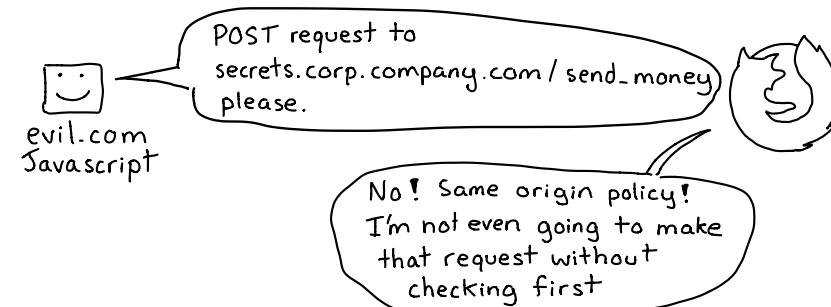
## Reason 1: cookies

Browsers often send your cookies with HTTP requests. You don't want evil.com to be able to make requests using your login cookies. They'd be logged in as you!



## Reason 2: network access

You might be on a private network (for example your company's corporate network) that evil.com doesn't have access to, but your computer does.



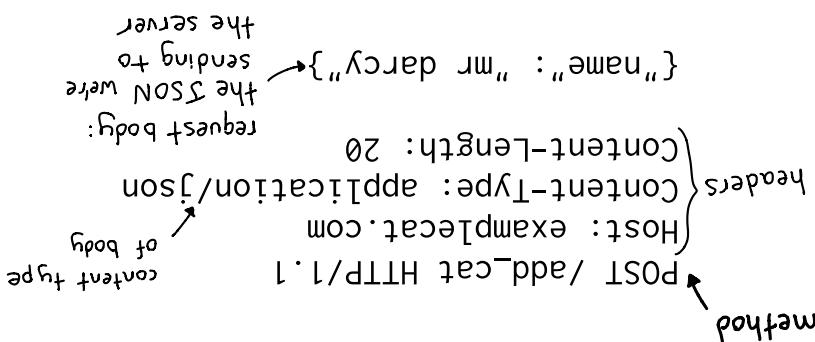
\* HTTP requests +  
anatomy of an

HTTP requests always have:

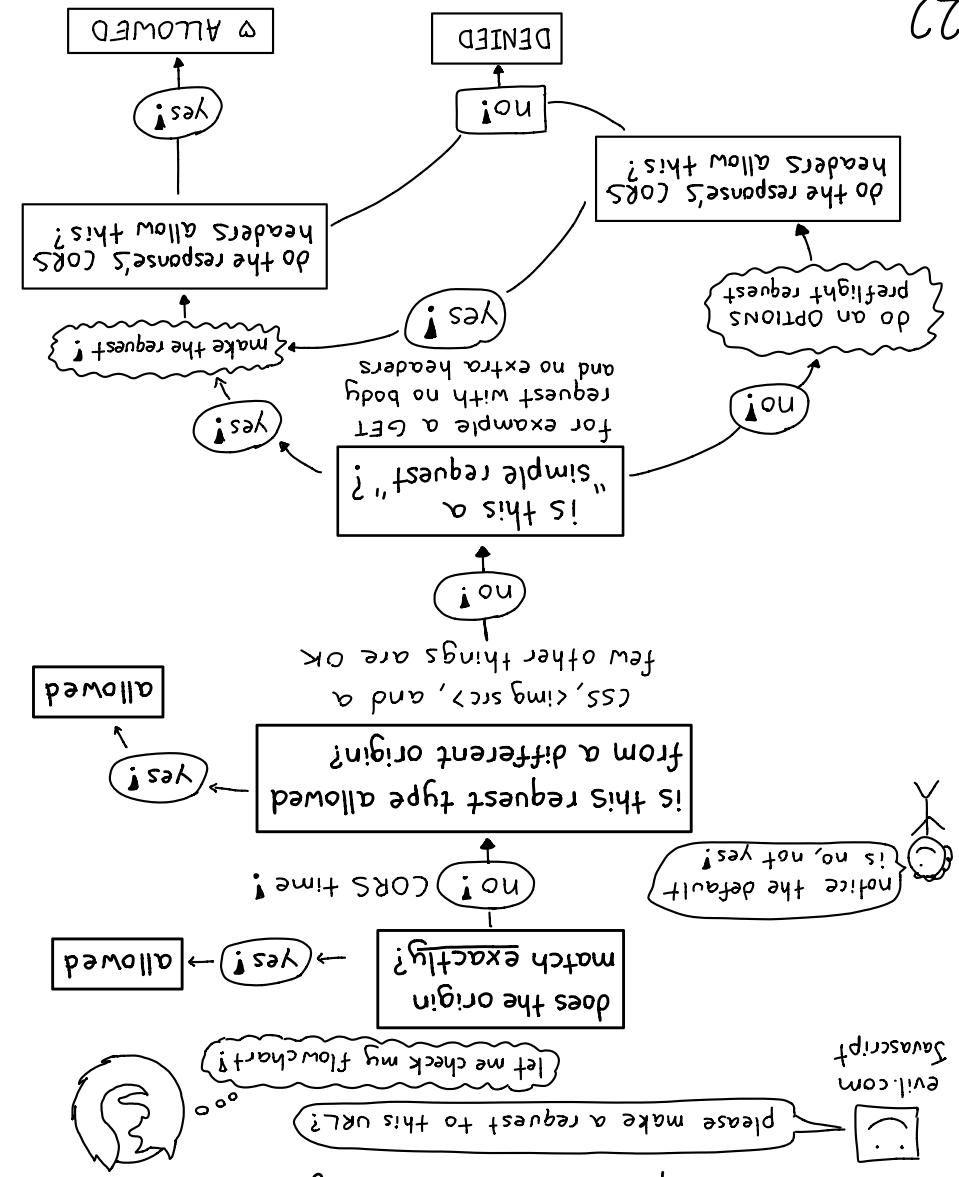
- a domain (like example.cat.com)
- a resource (like /cat.png)
- a method (GET, POST, or something else)
- headers (extra information for the server)
- There's an optional request body. GET requests usually don't have a body, and POST requests usually do.
- This is an HTTP 1.1 request for example.cat.com/cat.png
- It's a GET request, which is what happens when you type a URL in your browser. It doesn't have a body.

Usually GET or POST) resource being requested → method  
HTTP version  
GET /cat.png HTTP/1.1  
Host: example.cat.com  
User-Agent: Mozilla...  
headers  
Cookie: ....  
domain being requested

Here's an example POST request with a JSON body:



Here's an example POST request with a JSON body:



An origin is the protocol + domain including subdomains + port example: <https://babby.example.com:443> The same origin policy is one way browsers protect you from

the same original policy

# request methods

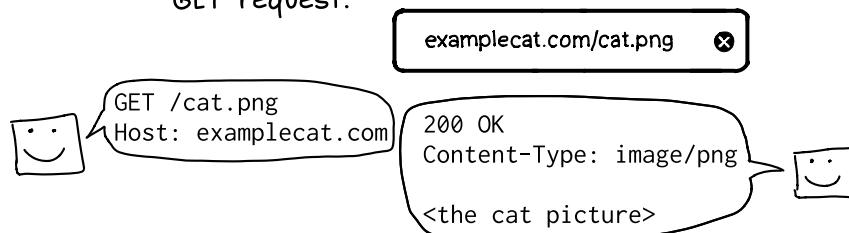
Every HTTP request has a method. It's the first thing in the first line:

→ this means it's a *GET* request

GET /cat.png HTTP/1.1

There are 9 methods in the HTTP standard. 80% of the time you'll only use 2 (GET and POST).

**GET** When you type an URL into your browser, that's a GET request.

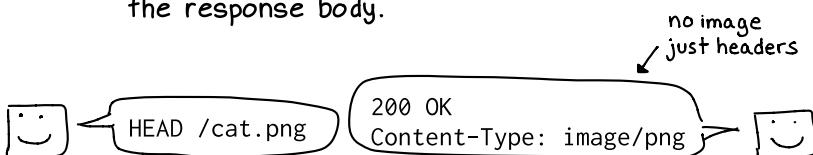


**POST** When you hit submit on a form, that's (usually) a POST request.



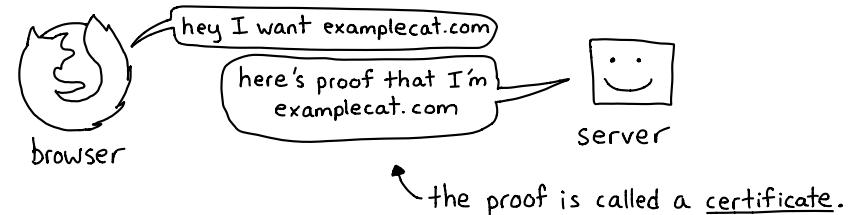
The big difference between GET and POST is that GETs are never supposed to change anything on the server.

**HEAD** Returns the same result as GET, but without the response body.



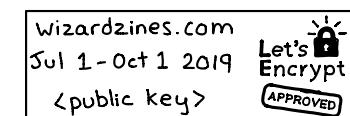
# certificates

To establish an HTTPS connection to examplecat.com, the client needs proof that the server actually is examplecat.com.



A TLS certificate has:

- a set of domains it's valid for (eg examplecat.com)
  - a start and end date (example: july 1 2019 to oct 1 2019)
  - a secret private key which only the server has
  - a public key to use when encrypting
  - a cryptographic signature from someone trusted
- this is the only secret part, the rest is public



The trusted entity that signs the certificate is called a ★ Certificate Authority ★ (CA) and they're responsible for only signing certificates for a domain for that domain's owner.



When your browser connects to examplecat.com, it validates the certificates using a list of trusted CAs installed on your computer. These CAs are called "root certificate authorities".

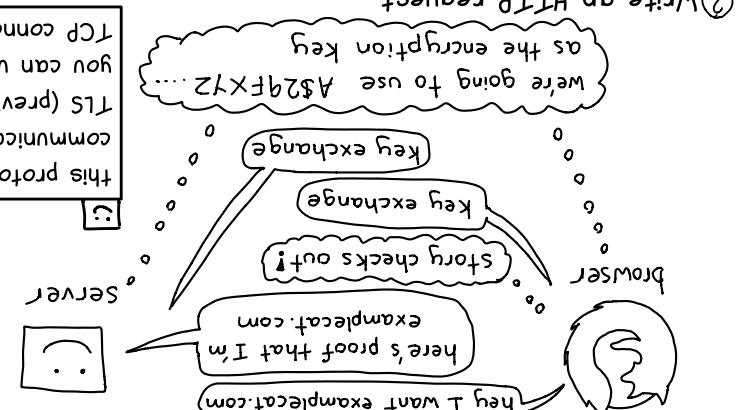


# HTTPS: HTTP + Secure ⚡

Here's what your browser does when it asks for `https://examplecat.com/cat.png`:

- Negotiate an encryption key (AES symmetric key) to use for this connection to `examplecat.com`. The browser and server will use the same key to encrypt/decrypt content.

Simplified version of how picking the encryption key works:



`https://examplecat.com/cat.png`:

- Negotiate an encryption key (AES symmetric key) to use for this connection to `examplecat.com`.
- Write an HTTP request: `GET /cat.png HTTP/1.1`
- Encrypt the HTTP request with AES & send it to `examplecat.com`:  
User-Agent: Mozilla/5.0 Host: examplecat.com GET /cat.png HTTP/1.1 ...  
...  
\$AFA9bbC4~833BF...
- Receive encrypted HTTP response:  
BXF v56 gxx ...  
...  
200 OK  
Content-Type: image/png  
nice, that means: ...

## OPTIONS

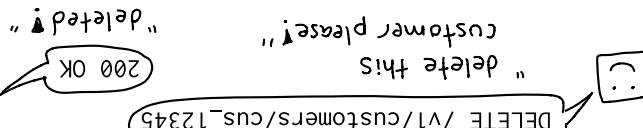
The CORS page has more about that.

It also tells you which methods are available.

OPTIONS is mostly used for CORS requests.

## DELETE

Used in many APIs (like the Stripe API) to delete resources.



## PUT

Used in some APIs (like the S3 API) to create or update resources. `PUT /cat/1234` lets you

`GET /cat/1234` later.

## PATCH

Used in some APIs for partial updates to a resource ("just change this field").

I've never seen a server that supports this, you probably don't need to know about it.

## TRACE

Different from all the others: instead of making a request to a server directly, it asks for a proxy to open a connection.

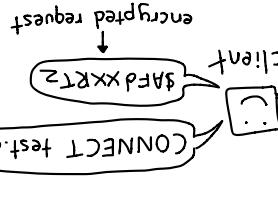
If you set the `HTTP_PROXY` environment variable to a proxy server, many HTTP libraries

will use this protocol to proxy your requests.

## CONNECT

Different from all the others: instead of making a request to a proxy server, many HTTP libraries

will use this protocol to proxy your requests.



## HTTP + Secure

# request headers

These are the most important request headers:

## Host

The domain.  
The only required header.

Host: examplecat.com User-Agent: curl 7.0.2 Referer: https://examplecat.com  
yes, it's misspelled!

## User-Agent

name + version of your browser and OS

## Referer

website that linked or included the resource

## Authorization

eg a password or API token  
base64 encoded user:password

Authorization: Basic YXZ

## Cookie

Send cookies the server sent earlier keeps you logged in.

Cookie: user=b0rk

## Range

lets you continue downloads ("get bytes 100-200")  
Range: bytes=100-200

## Cache-Control

"max-age = 60" means cached responses must be less than 60 seconds old

## If-Modified-Since

only send if resource was modified after this time

If-Modified-Since: Wed, 21 Oct...

## If-None-Match

only send if the ETag doesn't match those listed

If-None-Match: "e7ddac"

## Accept

MIME type you want the response to be

Accept: image/png

## Accept-Encoding

set this to "gzip" and you'll probably get a compressed response

Accept-Encoding: gzip

## Accept-Language

set this to "fr-CA" and you might get a response in French

Accept-Language: fr-CA

## Content-Type

MIME type of request body, e.g. "application/json"

## Content-Encoding

will be "gzip" if the request body is gzipped

## Connection

"close" or "keep-alive". Whether to keep the TCP connection open.

# HTTP/2

HTTP/2 is a new version of HTTP.

Here's what you need to know:

## ★ A lot isn't changing

All the methods, status codes, request/response bodies, and headers mean exactly the same thing in HTTP/2.

before (HTTP/1.1)

method: GET  
path: /cat.gif  
headers:

- Host: examplecat.com
- User-Agent: curl

after (HTTP/2)

method: GET  
path: /cat.gif  
headers:

- authority: examplecat.com
- User-Agent: curl

## ★ HTTP/2 is faster

Even though the data sent is the same, the way HTTP/2 sends it is different. The main differences are:

- It's a binary format (it's harder to tcpdump traffic and debug)
- Headers are compressed
- Multiple requests can be sent on the same connection at a time

before (HTTP/1.1)

→ request 1  
→ request 2  
response 1 ←  
response 2 ←

after (HTTP/2)

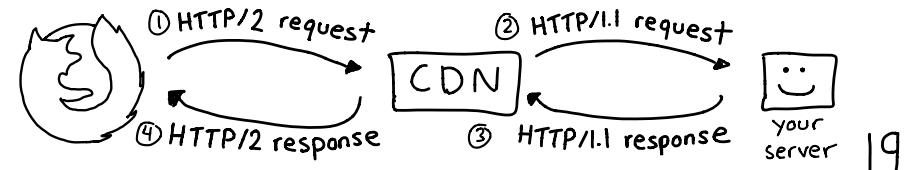
→ request 1  
→ request 2  
out of order { response 2 ←  
is OK { response 1 ← }

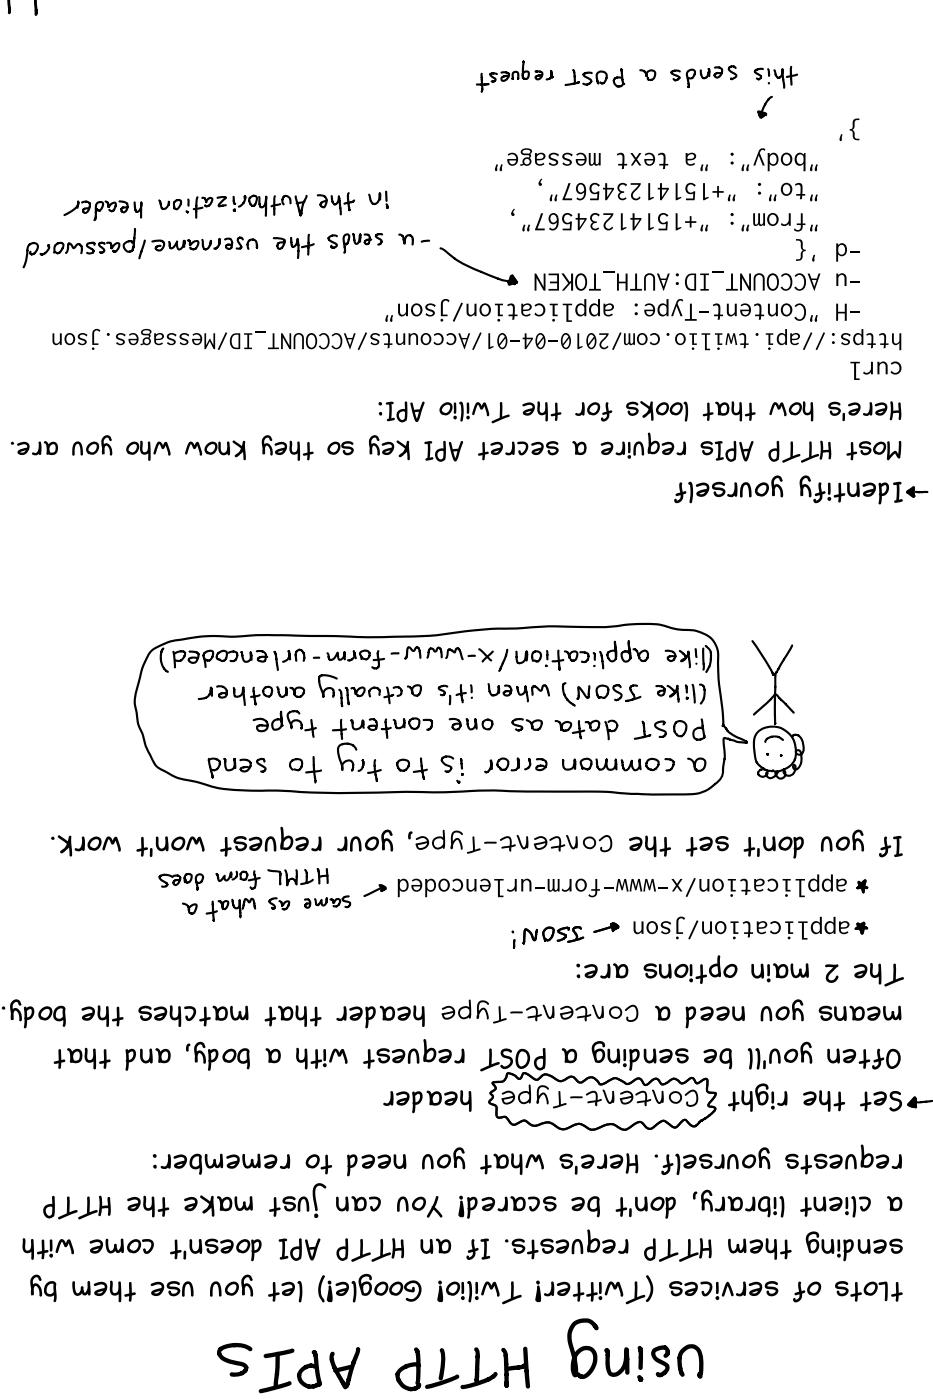
} one TCP connection

All these changes together mean that HTTP/2 requests often take less time than the same HTTP/1.1 requests.

## ★ Sometimes you can switch to it easily

A lot of software (CDNs, nginx) let clients connect with HTTP/2 even if your server still only supports HTTP/1.1.





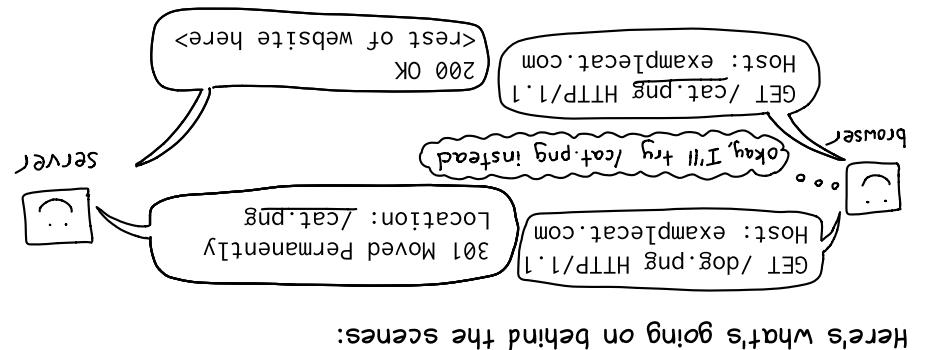
301 Moved Permanently redirects are PERMANENT: after a browser sees once, it'll always use examplecat.com/cat.png when someone types examplecat.com/dog.png forever. You can't take it back and decide to not to redirect. If you're not sure you want to redirect your site for eternity, use 302 Found to redirect instead.

Setting up redirects is a great thing to do if you move your site to a new domain!

The new URL doesn't have to be on the same domain: The Location header tells the browser what new URL to use.

The new URL doesn't have to be on the same domain: examplecat.com/panda can redirect to pandas.com.

Setting up redirects is a great thing to do if you move your site to a new domain!



but end up at a slightly different URL:  
examplecat.com/dog.png  
examplecat.com/cat.png  
I didn't type that!  
oh, where did the cat come from?  
I didn't type that!

Sometimes you type a URL into your browser:

## Redirections

# anatomy of an HTTP response

HTTP responses have:

- a status code (200 OK! 404 not found!)
- headers
- a body (HTML, an image, JSON, etc)

Here's the HTTP response from examplecat.com/cat.txt:

```

HTTP/1.1 200 OK status
Accept-Ranges: bytes
Cache-Control: public, max-age=0
Content-Length: 33
Content-Type: text/plain; charset=UTF-8
Date: Mon, 09 Sep 2019 01:57:35 GMT
Etag: "ac5affa59f554a1440043537ae973790-ssl"
Strict-Transport-Security: max-age=31536000
Age: 0
Server: Netlify

\  / cat! ^)
 ) ( ' ) ←
( / )
\(_)_/

```

} status code

} headers

} body

There are a few kinds of response headers:

when the resource was sent/modified:

Date: Mon, 09 Sep 2019 01:57:35 GMT  
Last-Modified: 3 Feb 2017 13:00:00 GMT

about the response body:

Content-Language: en-US Content-Type: text/plain; charset=UTF-8  
Content-Length: 33 Content-Encoding: gzip

caching:

ETag: "ac5affa..." Age: 255  
Vary: Accept-Encoding Cache-Control: public, max-age=0  
**security: (see page 25)**

X-Frame-Options: DENY Strict-Transport-Security: max-age=31536000  
X-XSS-Protection: 1 Content-Security-Policy: default-src https:

and more:

Connection: keep-alive Accept-Ranges: bytes  
Via: nginx  
Set-Cookie: cat=darcy; HttpOnly; expires=27-Feb-2020 13:18:57 GMT;

# caching headers

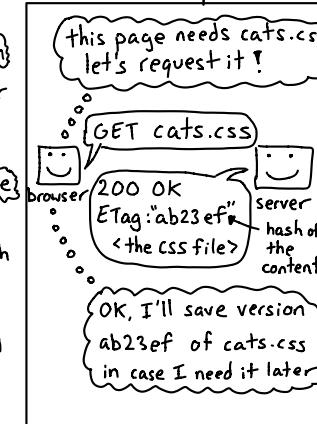
These 3 headers let the browser avoid downloading an unchanged file a second time.

initial request

**ETag**

response header  
and  
If-None-Match  
request header

If-Modified-Since  
is similar to  
If-None-Match  
but with  
Last-Modified  
instead of ETag



the next day



**Vary**  
response header

Sometimes the same URL can have multiple versions (spanish, compressed or not, etc).  
Caches categorize the versions by request header

like this:

Accept-Language	Accept-Encoding	content
en-US	-	hello
es-ES	-	hola
en-US	gzip	f\$xx99æf^.. (compressed gibberish)

The Vary header tells the cache which request headers should be the columns of this table.

**Cache-Control**

request AND response header

Used by both clients and servers to control caching behaviour. For example:

Cache-Control: max-age=999999999999  
from the server asks the CDN or browser to cache the thing for a long time.

Whether Range requests header  
is supported for this resource

### Access - Ranges

Whether body is compressed:  
Content-Encoding: gzip

### Content - Encoding

Length of body in bytes:  
Content-Length: 33

### Content - Length

Allow cross-origin requests:  
Access-Control-Allow-Origin:

### Access - Control - \*

"Close" or "keep-alive"  
Connection open:  
TCP connection open

### Connection

That response will  
be based on  
various headers

### Vary

When content was  
last modified  
(not always accurate)

### Last - Modified

Location: /cat.png  
URL to redirect to

### Location

Language of body:  
Content-Language: en-US

### Content - Language

MIME type of body:  
Content-Type: text/plain

### Content - Type

Sets a cookie:  
Set-Cookie: name=value; HttpOnly

### Set - Cookie

Added by proxy servers  
The response is stale  
and should be re-requested  
after this time.

### Via

Cache-control:  
max-age=300  
Cache-control: max-age=300

### Cache - Control

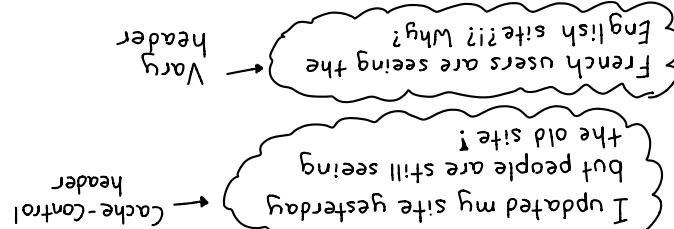
When response  
has been cached  
seconds response  
how many times  
Age

### Date

I paid  
for this?  
How will  
you owe me  
now you owe me  
\$100 for bandwidth  
<no response>

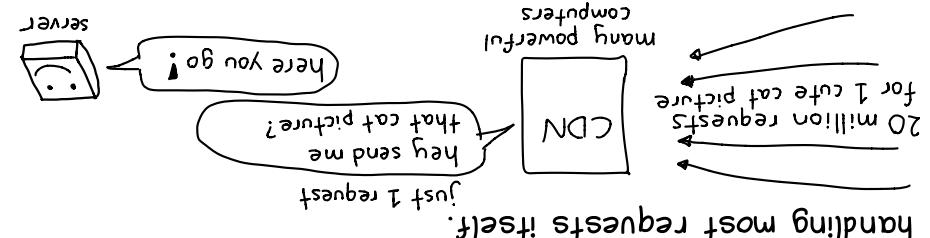
## response headers

Next, we'll explain the HTTP headers your CDN or browser uses to decide how to do caching.

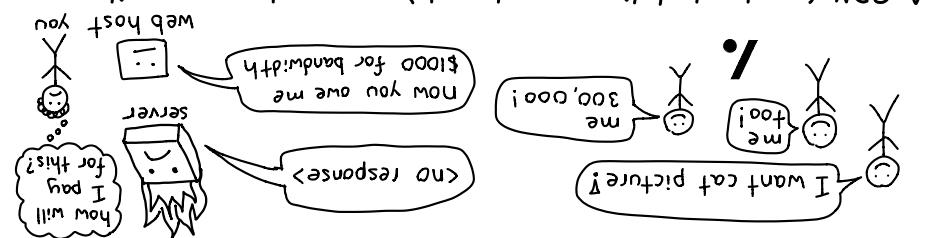


This is great but caching can cause problems too!

Today, there are many free or cheap CDN services available, which means if your site gets popular suddenly you can easily keep it running!



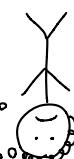
A CDN (content delivery network) can make your site faster and save you money by caching your site and handing most requests itself.



In 2004, if your website suddenly got popular, often the webserver wouldn't be able to handle all the requests.

## content delivery networks

browser uses to decide how to do caching.



...

# HTTP status codes

Every HTTP response has a ★status code★.



There are 50ish status codes but these are the most common ones in real life:

200 OK

301 Moved Permanently

302 Found  
temporary redirect

304 Not Modified

the client already has the latest version, "redirect" to that

400 Bad Request

403 Forbidden  
API key/OAuth/something needed

404 Not Found  
we all know this one :)

429 Too Many Requests  
you're being rate limited

500 Internal Server Error  
the server code has an error

503 Service Unavailable  
could mean nginx (or whatever proxy)  
couldn't connect to the server

504 Gateway Timeout  
the server was too slow to respond

} 2xx's mean  
★ success ★

} 3xx's aren't errors, just redirects to somewhere else

} 4xx errors are generally the client's fault: it made some kind of invalid request

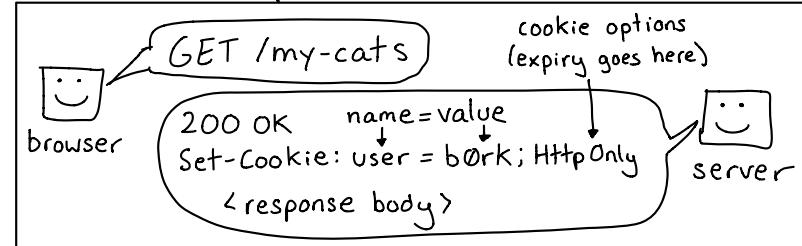
} 5xx errors generally mean something's wrong with the server.

# how cookies work

Cookies are a way for a server to store a little bit of information in your browser.

They're set with the Set-Cookie response header, like this:

first request: server sets a cookie



Every request after: browser sends the cookie back



Cookies are used by many websites to keep you logged in. Instead of user=b0rk they'll set a cookie like sessionid=long-incomprehensible-id. This is important because if they just set a simple cookie like user=b0rk, anyone could pretend to be b0rk by setting that cookie!

Designing a secure login system with cookies is quite difficult — to learn more about it, google "OWASP Session Management Cheat Sheet".