

love this?

find more awesome zines at

→ [jvns.ca/zines](http://jvns.ca/zines) ←

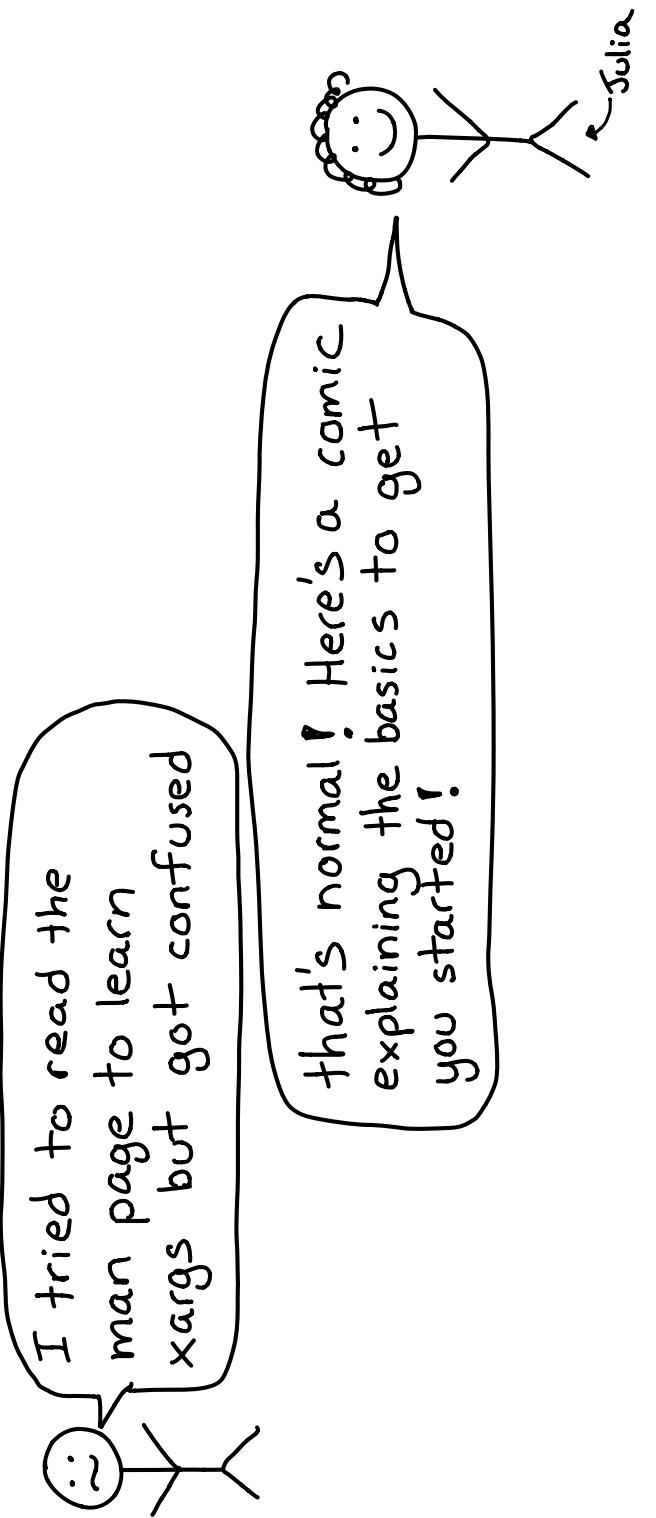
CC-BY-NC-SA computer wizard industries 2018

# RiteSize Command Line

By Julia Evans



This zine explains some of the most useful Unix command line tools in 1 page each.



Even if you've used the command before, I might have a new trick or two for you ❤

\*

## more useful tools

- make
- screen
- diff -U
- jq
- tmux
- wipe
- nohup
- date
- entr
- fish
- disown
- ranger
- seq
- join
- chronic
- xxd
- parallel:
  - GNU parallel
- hexdump
- objdump
- pigz / pixz
- sort --parallel
- strings

\*

# lsof

## lsof

stands for list open files

what lsof tells you  
for each open file:

- pid
- file type (regular? directory?  
FIFO? socket?)
- file descriptor (FD column)
- user
- filename / socket address

**-P PID**  
list the files PID has  
open

**lsof /some /dir**

list just the open files  
in /some /dir

**-i**  
list open network sockets  
(sockets are files!)

examples:

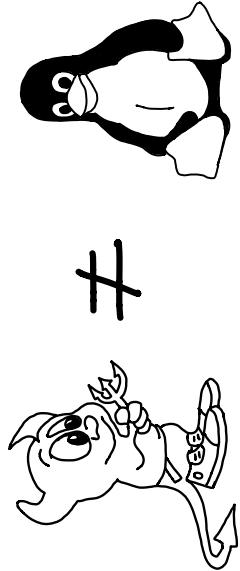
**-i -n -p** ← -n & -p mean  
"don't resolve  
-i :8080 host names / ports"  
-i TCP (also -Pn);  
-i -s TCP:LISTEN

**find deleted files**  
**\$ lsof | grep deleted**  
will show you deleted files!  
You can recover open deleted  
files from  
**/proc/<pid>/fd/<fd>**  
process that opened the file

**netstat**  
another way to list open  
sockets on Linux is:  
**netstat -tunapl**  
**tuna, please!**  
On Mac netstat has  
different args.

BSD#GNU.....4	bash tricks.....10-11	misc commands...17
find.....6	tar.....13	head & tail.....18
grep.....5	disk usage.....12	less.....19
xargs.....7	ps.....14	kill .....
awk.....8	top.....15	cat.....21
sed.....9	sort & uniq.....16	lsof.....22

## Table of contents

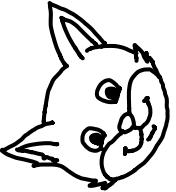


For almost all these tools, there are at least 2 versions:

① The BSD version (on BSDs & Mac OS)

② The GNU version (on Linux)

All of the examples in this zine were tested on Linux. Some things (like sed -i) are different on Mac. Be careful when writing cross-platform scripts! You can install the GNU versions on Mac with `brew install coreutils`.



## cat & friends

21

**cat -n**  
prints out the file with line numbers!

- 1 Once upon a midnight...
- 2 Over many a quaint...
- 3 While I nodded, nearly

You can use cat as an EXTREMELY BASIC text editor:

- ① Run \$ cat > file.txt
- ② type the contents (don't make mistakes !)
- ③ press ctrl+d to finish

**cat** concatenates files

```
$ cat myfile.txt
prints contents of myfile.txt
$ cat *.txt
prints all .txt files put together!
```

### zcat

cats a gzipped file!  
Actually just a 1-line shell script that runs `gzip -cd`, but easier to remember.



### tee

'tee file.txt' will write its stdin to both stdout and file.txt  
→ stdout  
→ file.txt  
stdin → tee a.txt → a.txt

how to redirect to a file owned by root  
**\$ sudo echo "hi" >> x.txt**

/ this will open x.txt as your user, not as root, so it fails!  
**\$ echo "hi" | sudo tee -a x.txt**

will open x.txt as root ↴



# find

6

**find** searches a directory for files  
 find /tmp -type d, -print  
 directory which files action to do with the files  
 to search

**-size 0**  
 find empty files!  
 Useful to find files you created by accident

**-name /-iname**  
<sup>case insensitive</sup>  
 the filename! eg  
 -name '\*.txt'

**-path /-ipath**  
 search the full path!  
 -path '/home/\*/\*.go'

**-print0**  
 print null-separated filenames  
 Use with xargs -0!

**-delete**  
 action: delete all files found

**-exec COMMAND**  
 action: run COMMAND on every file found

**-type [TYPE]**  
 f: regular file      l: symlink  
 d: directory      + more!  
**-maxdepth NUM**  
 only descend NUM levels when searching a directory

**locate**  
 The locate command searches a database of every file on your system.  
 good: faster than find  
 bad: can get out of date  
**\$sudo updatedb**  
 updates the database

# less

19

**less is a pager**  
 that means it lets you view (not edit) text files or piped in text  
 man uses your pager (usually less) to display man pages

many vim shortcuts work in less  
 / search  
 n/N next / prev match  
 j/JK down/up a line  
 m/m mark/return to line  
 g/G beginning/end of file  
 † (gg in vim)

**q** quit !!  
**V** ~ lowercase edit file in your \$EDITOR  
 arrow keys, Home/End, Pg Up, Pg Dn work in less

**less -r**  
 displays bash escape codes as colours  
 try ls --color | less -r  
 without -r  
 a.txt      a.txt+  
 ESC[10m ESC[10;31m a.txt q2      ESC[0m  
**a.txt.gz**      red,bold

**+**  
 + runs a command when less starts  
 less **+F** : follow updates  
 less **+G** : start at end of file  
 less **+20%** : start 20% into file  
 less **+/foo** : search for 'foo' right away

# head & tail

## head

shows you the first 10 lines of a file

if you pipe a program's output to head, the program will stop after printing 10 lines (it gets sent `SIGPIPE`)

## -c NUM

show the first /last NUM bytes of the file

`head -c 1k`

will show the first 1024 bytes

## tail

tail shows the last 10 lines!

`tail -f FILE` will follow:

print any new lines added to the end of FILE. Super useful for log files!

## tail --retry

keep trying to open file if its inaccessible

## tail --pid PID

stop when process PID stops running (with -f)

## -n NUM

-n NUM (either head or tail) will change the # lines shown

`head -n -NUM` } show all tail -n +NUM } but the last / first NUM lines

## tail --follow=name

Usually tail -f will follow a file descriptor.

`tail --follow=name FILENAME` will keep following the same file name, eg if the file descriptor is rotated.

# Xargs

this is useful when you want to run the same command on a list of files!

- delete (`xargs rm`)
- combine (`xargs cat`)
- search (`xargs grep`)
- replace (`xargs sed`)

how to replace "foo" with "bar" in all .txt files:

```
find . -name '*.txt' |  
xargs sed -i $'foo/bar/g'
```

how to lint every Python file in your Git repo:

```
git ls-files | grep .py |  
xargs pep8
```

more useful xargs options

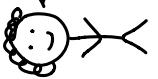
`-n 1` makes xargs run a separate process for each input.  
`-P` is the max number of parallel processes xargs will start

# awk

8

awk is a tiny programming language for manipulating columns of data

I only know how to do 2 things with awk but it's still useful!



basic awk program structure

```
BEGIN{ ... }  
      CONDITION {action}  
      CONDITION {action}  
END{...}    do action on  
           lines matching  
           condition
```

so being able to get the column you want with awk is GREAT

so many unix commands print columns of text (ps! ls!)

awk program example:  
sum the numbers in the 3rd column

```
action  
{ s += $3 };  
END {print s}  
at the end, print  
the sum!
```

extract a column of text with awk

awk -F, ' {print \$5}'

single quotes! 5<sup>th</sup> column separator

(this is 99% of what I do with awk)



awk program example:  
print every line over 80 characters

length(\$0) > 80

(there's an implicit {print} as the action)

# misc commands

pv

"pipe viewer", gives you stats on data going through a pipe

cal

a tiny calendar ☺

ncdu

figure out what kind all your disk space

comm

format input into columns

xsel / xclip

copy/paste from system clipboard.  
(pbcopy / pbpaste on Mac)

diff

diff 2 files. Run with -U 8' for context.

+s

add a timestamp in front of every input line

# sort & uniq

**sort** sorts its input

\$ sort names.txt

the default sort is alphabetical.

**uniq** removes duplicates

a  
b  
b => a  
a  
c  
c  
notice there  
are still 2  
'a's!  
only uniqueness  
adjacent  
matching lines

Sort -n	
numeric sort	'sort' order
'sort' order	'sort-n' order
12	12
15000	=
48	48
96	96
6020	6020
96	15000

sort + uniq = ❤

Pipe something to 'sort | uniq' and you'll get a deduplicated list of lines! **sort -u** does the same thing.

b  
b  
b  
| sort -u => a  
or sort | uniq

**sort -h**: human sort

'sort-n' order

'sort-h' order

15 G	45 K
30 M	30 M
=	=
45 K	15 G
~ 200G	~ 200G

useful example:  
du -sh \* | sort -h

**uniq -c**

counts each line it saw.

Recipe: get the top 10 most common lines in a file:

\$ sort foo.txt  
| uniq -c  
| sort -n  
| tail -n 10

I use this a lot

# Sed

q

some more sed incantations...

**sed -n 12 p**

print 12<sup>th</sup> line

**sed s+cat/+dog/+**  
use + as a regex delimiter  
way easier than escaping is like s/cat\/l/dog\//!

**sed -n s/cat/dog/p**  
only print changed lines

**sed 6**  
double space a file (good for long error lines)

**sed 'cat/a dog'**  
append 'dog' after lines containing 'cat'

**sed -n 5,30 p**  
print lines 5-30

# bash tricks

10

## \* **ctrl + r \***

search your history!

I use this ❤ constantly ❤ to rerun commands

```
$ convert file.jpg file.png  
$ convert file.jpg file.png  
{1..5} expands to 1 2 3 4 5  
(for i in {1..100};)
```

## \* **magical braces \***

```
$ convert file.{jpg,png}  
expands to
```

```
$ convert file.jpg file.png  
{1..5} expands to 1 2 3 4 5  
(for i in {1..100};)
```

## loops

```
for i in *.png  
do  
convert $i $i.jpg  
done
```

for loops:  
easy & useful!

## \$()

gives the output of a command

```
$ touch file-$(date -I)  
↑  
create a file named  
file-2018-05-25
```

# top

## top

a live-updating summary of the top users of your system's resources

who's using all my memory  
top  
chromium  
obv

let's explain some numbers in top!

## load average

3 numbers that roughly reflect demand for your CPUs on the system in the last 1, 5, and 15 minutes if it's higher than the # of CPUs you have, that's often bad

15

## memory

4 numbers:

total / free / used / cached  
One perhaps unexpected thing:  
total is **not** free + used!  
total = free + used + cached  
filesystem cache

15

## RES

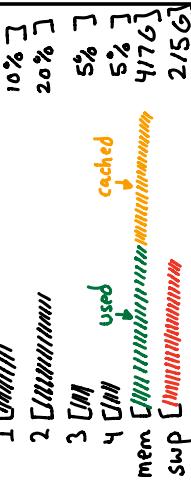
this column is the "resident set size" aka how much RAM your process is using.

SHR  
RES is how much of the RES is shared with other processes

## % CPU

what?

this column is given as % of a single core. If you have 4 cores, this can go up to 400 %!



# PS

## ps

ps shows which processes are running

I usually run ps like this:

`$ ps aux`

u means include  
username column  
(ps -ef works too)

**process state**

Here's what the letters in ps's STATE column mean:

R: running  
S/D: asleep  
Z: zombie  
I: multithreaded  
+: in the foreground

## w

is for wide. ps auxww will show all the command line args for each process

## e

is for environment. ps auxe will show the environment vars!

**f** is for "forest" !. ps auxf will show you an ASCII art process tree!

**pstree** can display a process tree too

you can choose which columns to show with ps (ps -eo ...) One cool column is 'wchan' which tells you the name of the kernel function if the process is sleeping  
**try it:**  
`ps -eo user,pid,wchan,cmd`

## wchan

ps has 3 different sets of command line arguments ↪  
1. UNIX ( 1 dash)  
2. BSD (no dash)  
3. GNU ( 2 dashes)  
You can write monstrously like:  
`$ ps f -f ↑ forest(BSD) full format (UNIX)`

# more bash tricks 11

## cd -

changes to the directory you were last in

pushd & popd let you keep a stack

## <(< )

process substitution treat process output like a file (no more temp files!)

eg:  
`$ diff <(ls) <(ls -a)`

## ctrl+z

suspends (SIGTSTP) the running program

## fg

brings backgrounded/suspended program to the foreground

## bg

starts suspended program & backgrounds it (use after ctrl+z)

## shellcheck

shell script linter! helps spot common mistakes.

## type

tells you if something is a builtin, program, or alias try running type on `{time} {ping}` (pushd) (they're all different types)

# disk usage

12

## du

tells you how much disk space files / directories take up

**-S** summary: total size of all files in a directory

**-h** human readable sizes

## df

\* **df** \* tells you how much free space each partition has.

for human-readable sizes

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/sda3	18G	6	2.5G	86%	/
udev	4.83M	4.0K	4.83M	1%	/dev
tmpfs	99M	1.4M	97M	2%	/run
/dev/sda4	167G	157G	9.9G	95%	/home

## df -i

instead of % disk free, report how many **inodes** are used / free on each partition

running out of inodes is VERY ANNOYING - You can't create new files!

## ncdu

see what's using disk space navigate with arrow keys

17.5 GiB	[#####]	/music
3.2 GiB	[##	/pictures
5.7 MiB	] /text	
2.0 MiB	] file.pdf	

## iostat

get statistics about disk reads / writes

interval to report at  
# iostat 5

Device: kB\_read/s kB\_wrtn/s

sda 2190.21 652.87

sdb 6.00 0.00

13

# tar

The tar file format combines many files into one file.

a.txt dir/c.txt  
b.txt

.tar files aren't compressed by themselves. Usually you gzip them: .tar.gz or .tgz!

-X is for extract into the current directory by default (change with -C)

-C is for create makes a new tar file!

-t is for list lists the contents of a tar archive

-f is for file which tar file to create or unpack & more! see the man page :)

putting it together

**list** contents of a .tar.b22:

\$ tar -tjvf file.tar.b22  
\$ tar -c2f file.tar.gz dir/  
files to go in the archive

-Z gzip format (.gz)

-j bzip2 format (.bz2)

-J xz format (.xz)