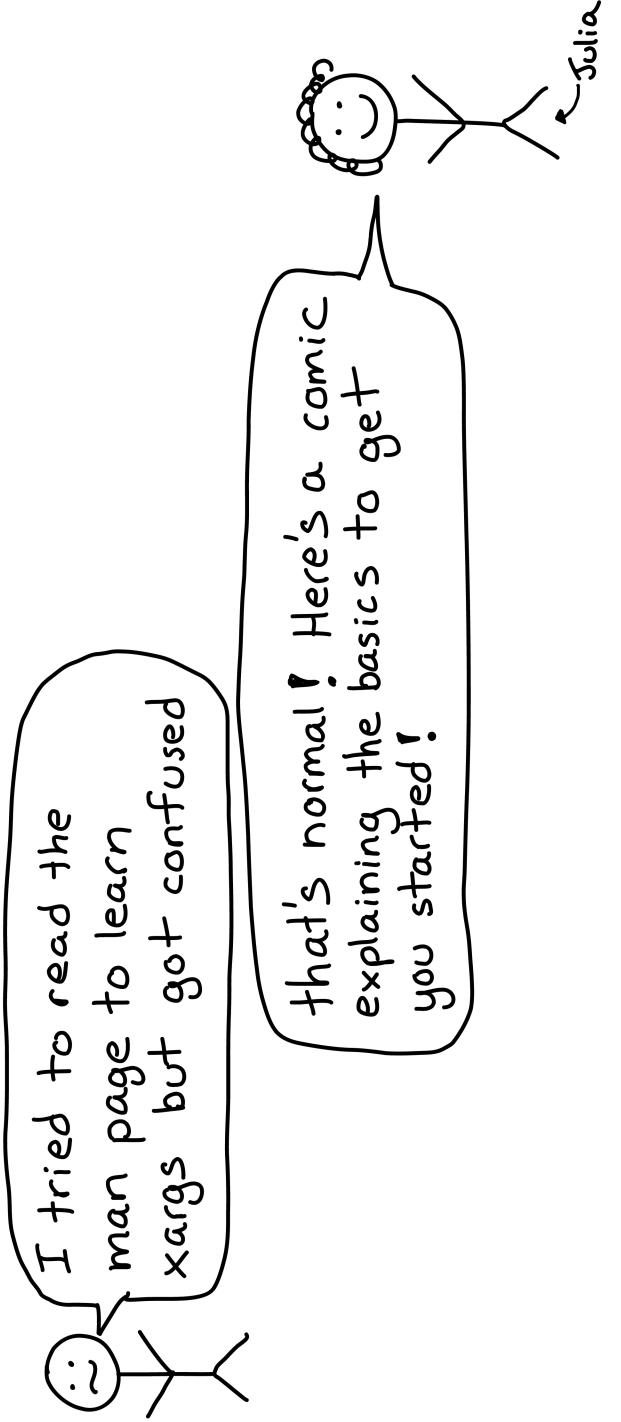


CC-BY-NC-SA computer wizard industries 2018

love this?
find more awesome zines at

This zine explains some of the most useful Unix command line tools in 1 page each.



Even if you've used the command before, I might have a new trick or two for you ❤

*

more useful tools

- make
- jq
- nohup
- disown
- cut/paste
- sponge
- xxd
- hexdump
- objdump
- strings
- screen
- tmux
- date
- entr
- seq
- join
- parallel:
 - GNU parallel
 - pigg / pixz
 - sort --parallel
- diff -U
- vipe
- imagemagick
- fish
- ranger
- chronic

*

lsof

lsof

stands for list open files



what lsof tells you
for each open file:

- pid
- file type (regular? directory?
FIFO? socket?)
- file descriptor (FD column)
- user
- filename / socket address

-P PID
list the files PID has
open

lsof /some /dir
list just the open files
in /some /dir

-i

list open network sockets
(sockets are files!)

examples:

```
-i -n -P ← -n & -P mean "don't resolve"
-i :8080      host names / ports"
-i TCP        (also -Pn)
-i -s TCP:LISTEN
```

find deleted files

\$ lsof | grep deleted

will show you deleted files!

You can recover open deleted
files from
`/proc/<pid>/fd/<fd>`

process that opened the file

netstat

another way to list open
sockets on Linux is:

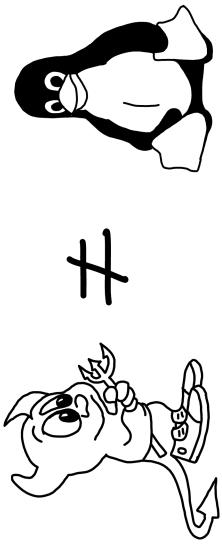
netstat -tunapl

↑
tuna, please!

On Mac netstat has
different args.

♥ Table of contents ♥

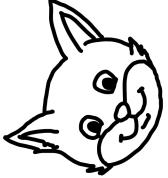
BSD#GNU	4	bash tricks.....	10-11	misc commands...17	
grep	5	disk usage.....	12	head & tail.....18	
find	6	tar	13	less.....19	
xargs.....	7	ps	14	kill	20
awk	8	top.....	15	cat.....21	
sed	9	sort & uniq.....	16	lsof.....22	



For almost all these tools, there are at least 2 versions:

- ① The BSD version (on BSDs & macOS)
 - ② The GNU version (on Linux)

All of the examples in this zine were tested on Linux.
Some things (like sed -i) are different on Mac.
Be careful when writing cross-platform scripts!
You can install the GNU versions on Mac with
`brew install coreutils`.



cat & friends

```
$ cat myfile.txt  
prints contents of myfile.txt  
  
$ cat *.txt  
prints all .txt files put  
together!
```

you can use cat as an
EXTREMELY BASIC text
editor:

- ① Run \$ cat > file.txt
 - ② type the contents (don't make mistakes ↴)
 - ③ press ctrl+d to finish

`cat -n` prints out the file with line numbers!

- 1 Once upon a midnight...
- 2 Over many a quaint.
- 3 While I nodded, nearly

zcat
cats a gzipped file!
Actually just a 1-line
shell script that runs
'gzip -cd', but easier
to remember.

tee

'tee file.txt' will write its stdin to both stdout and file.txt

find

6

-name / -iname
case insensitive
the filename! eg
-name '* .txt'

-type [TYPE]
f: regular file l: symlink
d: directory + more!

find searches a directory for files
find /tmp -type d -print
directory which files action to do with the files
to search

-path / -i-path
search the full path!
-path '/home /*/* .go'

-size 0

find empty files!
Useful to find files you created by accident

-exec COMMAND

action: run COMMAND on every file found

19

less

many vim shortcuts work in less

- / search
- n/N next / prev match
- j/k down / up a line
- m/ mark / return to line
- g/G beginning / end of file
- (gg in vim)

less is a pager
that means it lets you view (not edit) text files or piped in text man uses your pager (usually less) to display man pages

+

+ runs a command when less starts
less +F : follow updates
less +G : start at end of file
less +20% : start 20% into file
less +/foo : search for 'foo' right away

F

press F to keep reading from the file as it's updated (like tail -f)
press Ctrl+C to stop reading updates

quit !!

V lowercase edit file in your \$EDITOR arrow keys, Home / End, Pg Up, Pg Dn work in less

head & tail

18

head

shows you the first 10 lines of a file

if you pipe a program's output to head, the program will stop after printing 10 lines (it gets sent SIGPIPE)

-c NUM

show the first/last NUM bytes of the file

head -c 1K
will show the first 1024 bytes

tail

tail shows the last 10 lines!

tail -f FILE will follow:
print any new lines added to the end of FILE. Super useful for log files!

tail --retry

keep trying to open file if it's inaccessible

tail --pid PID
stop when process PID stops running (with -f)

-n NUM

-n NUM (either head or tail)
will change the # lines shown
 $\{ \text{head } -n \text{ -NUM} \text{, tail } -n \text{ +NUM} \}$ show all last / first NUM lines

tail --follow=name

Usually tail -f will follow a file descriptor.

tail --follow=name FILENAME
will keep following the same filename, eg if the file descriptor is rotated.

Xargs

7

how to replace "foo" with "bar" in all .txt files:

find . -name '* .txt' |
xargs sed -i s/foob/ar/g
→ search (xargs grep)
→ replace (xargs sed)

xargs takes whitespace separated strings from stdin and converts them into command-line arguments

```
$ echo "/home /tmp" | xargs ls  
will run  
ls /home /tmp
```

how to lint every Python file in your Git repo:

```
git ls-files | grep .py |  
xargs pep8
```

if there are spaces in your filenames "my day.txt"

xargs will think it's 2 files "my" and "day.txt"

fix it like this:

```
find . -print0 |  
xargs -0 COMMAND
```

more useful xargs options

(-n 1) makes xargs run max-args for each input.
capital P is the max number of parallel processes xargs will start

awk

8

awk is a tiny programming language for manipulating columns of data

I only know how to do 2 things with awk but it's still useful!

basic awk program structure

```
BEGIN{ ... }  
CONDITION {action}  
CONDITION {action}  
END {....} do action on  
lines matching  
CONDITION
```

so many unix commands print columns of text (ps! ls!) so being able to get the column you want with awk is GREAT

awk program example:
sum the numbers in
the 3rd column
action

```
{ s += $3 };  
END {print s};  
at the end, print  
the sum!
```

extract a column of text with awk

```
awk -F, ' {print $5}'  
column separator quotes! 5th column  
this is 99% of what I do with awk
```

awk program example:
print every line over
80 characters

```
[length($0) > 80]  
condition  
(there's an implicit  
{print} as the action)
```

misc commands ❤ 17

pv
"pipe viewer", gives you stats on data going through a pipe

cml
a tiny calendar ☰

diff
diff 2 files. Run with '-u 8' for context.

ncdu
figure out what kind of file (png? pdf?) a file is

xsel / xclip
copy/paste from system clipboard.
(pbcopy / pbpaste on Mac)

rlwrap
adds history & ctrl support to REPLs that don't already have them (rl stands for readline)

```
$ rlwrap python
```

ts
add a timestamp in front of every input line

comm
find lines 2 sorted files have in common

Sort & uniq

sort sorts its input

```
$ sort names.txt
the default sort is alphabetical.
```

uniq removes duplicates

a
b
b
a
a
c
c
notice there
are still 2
'a's! uniq
only uniquifies
adjacent
matching lines

sort -n
numeric sort
'sort' order
'sort-n' order

=	12	=	12
=	15000	=	48
=	48	=	96
=	6020	=	6020

15000 45 K 30 M 156 2006 2006

sort -h: human sort
'sort-n' order
'sort-h' order

156	45 K
30M	30M
45K	156
2006	2006

useful example:
du -sh * | sort -h

sort + uniq = ❤
Pipe something to
'sort | uniq', and you'll
get a deduplicated list
of lines! {sort-w} does the
same thing.
b
a
b
a
| sort -u => a
b
a
or sort | uniq

uniq -c
counts each line it saw.
Recipe: get the top 10 most
common lines in a file:
\$ sort foo.txt
| uniq -c
| sort -n
| tail -n 10

I use this a lot

Sed

q

Some more sed incantations...

sed -n 12 p

print 12th line

-n suppresses output so
only what you print with 'p'
gets printed

sed 5d
delete 5th line

sed /cat/d

delete lines matching /cat/

sed -n 5,30 p
print lines 5-30

change a file in place
with -i
in GNU sed it's -i
in BSD sed, -i SUFFIX
confuses me every time.

sed s+cat/+dog/+
use + as a regex delimiter
way easier than
escaping / is like
s/cat//dog//!

sed s+cat/+dog/+
double space a file
(good for long error lines)

sed '/cat/a dog'
append 'dog' after lines
containing 'cat'

sed 'i7 panda'
insert "panda" on line 17

bash tricks

10

* ctr + r *

search your history!

I use this constantly ❤️ to rerun commands

* magical braces *

```
$ convert file.{jpg/png}
```

```
$ convert file.jpg file.png  
{1..5} expands to 1 2 3 4 5  
(for i in {1..100}; ...)
```

School

```
for i in *.png  
do  
    convert $i $i.jpg  
done
```

for loops:
easy & useful!

§ ()

gives the output of a command

```
$ touch file-$(date -I)  
↑  
create a file named  
file-2018-05-25
```

more keyboard shortcuts

- ctrl+a beginning of line
- ctrl+e end of line
- ctrl+l clear the screen
- + lots more emacs
shortcuts too !

15

• load average too

top

a live-updating summary of
the top users of your
system's resources

who's using
all my
memory

chrome to
obu

top

load average

3 numbers that roughly reflect demand for your CPUs on the system in the last 1, 5, and 15 minutes if it's higher than the # of CPUs you have, that's often bad

- memory
- 4 numbers:
 - total / free / used / cached
 - One perhaps unexpected thing:
total is not free + used!
 - total = free + used + cached

memory

- 4 numbers:
 - total / free / used / cached
 - One perhaps unexpected thing
 - Total is not free + used!
 - Total = free + used + cached filesystem cache

% CPU

what? 350%? 00000

This column is given as % of a single core. If you have 4 cores, this can go up to 400%.

RES

this column is the "resident set size" aka how much RAM your process is using.

SHR is how much of the RES is shared with other processes

htod

a prettier & more interactive
version of top *

PS

ps

ps shows which processes are running

I usually run ps like this:

```
$ ps aux
a+u  
u means include + together x show all * process
username column
(ps -ef works too)
```

process state

Here's what the letters in ps's STATE column mean:

- R: running
- S/D: asleep
- Z: zombie
- t: multithreaded
- +: in the foreground

e

is for environment. ps auxe will show the environment vars!

wchan
you can choose which columns to show with ps (ps -eo ...) One cool column is 'wchan' which tells you the name of the kernel function if the process is sleeping
try it:
ps -eo user,pid,wchan,cmd

ps has 3 different sets of command line arguments ↪

1. UNIX (1 dash)
2. BSD (no dash)
3. GNU (2 dashes)

you can write monstrously like:
\$ ps f -f
↑
forever(BSD) → full format (UNIX)

more bash tricks

11

cd -

changes to the directory you were last in

```
pushd & popd let you keep a stack
```

<()

process substitution
treat process output like a file (no more temp files!)

```
eg: $ diff <(ls) <(ls -a)
```

ctrl+z

suspends (SIGTSTP) the running program

fg

brings background/suspended program to the foreground

bg

starts suspended program & backgrounds it (use after ctrl+z)

shellcheck

shell script linter! helps spot common mistakes.

type

tells you if something is a builtin, program, or alias
try running type on {time ping pushd} (they're all different types)

disk usage

12

du

tells you how much disk space files / directories take up

`-s` summary: total size of all files in a directory

`-h` human readable sizes

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/sda3	18G	6	2.5G	86%	/
udev	483M	4.0K	483M	1%	/dev
tmpfs	99M	1.4M	97M	2%	/run
/dev/sda4	167G	157G	9.9G	95%	/home

df -i

instead of % disk free, report how many inodes are used/ free on each partition

running out of inodes is VERY ANNOYING - you can't create new files!

ncdu

see what's using disk space navigate with arrow keys

17.5 GiB	[#*##*#]	/music
3.2 GiB	[##]	/pictures
5.7 MiB	[]	/text
2.0 MiB	[]	/file.pdf

iostat

get statistics about disk reads / writes interval to report at

```
# iostat 5
Device: kB_read/s kB_wrtn/s
sda      2190.21   652.87
sdb       6.00     0.00
```

tar

13

The .tar file format combines many files into one file.

a.txt
b.txt
c.txt
d.txt

.tar files aren't compressed by themselves. Usually you gzip them: .tar.gz or .tgz!

Usually, when you use the 'tar' command, you'll run some incantation To unpack a tar.gz, use:

`-tar -xzf file.tar.gz`

let's learn!

what's xzf?

`-t` is for list lists the contents of a tar archive

`-f` is for file which tar file to create or unpack

tar can compress / decompress

- `-Z` gzip format (.gz)
- `-j` bzip2 format (.bz2)
- `-J` xz format (.xz)

putting it together

- `list` contents of a .tar.bz2:
`$ tar -tzvf file.tar.bz2`
- `create` a .tar.gz:
`$ tar -czf file.tar.gz dir/`
- `files to go` in the archive