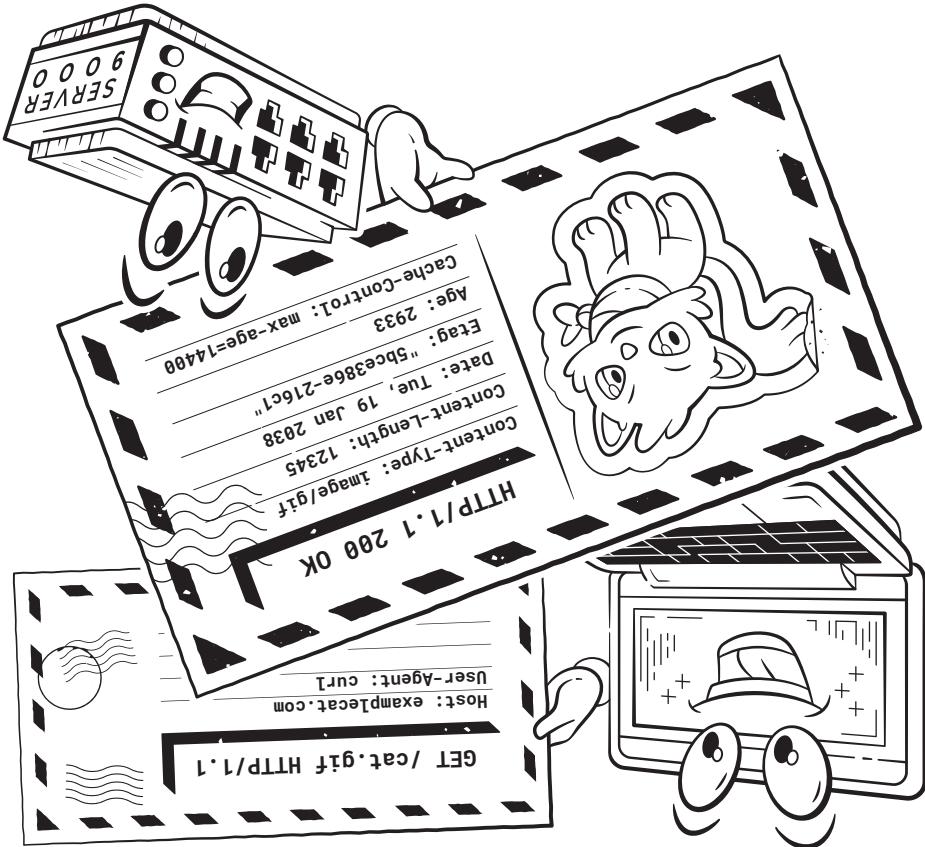


by Julia Evans

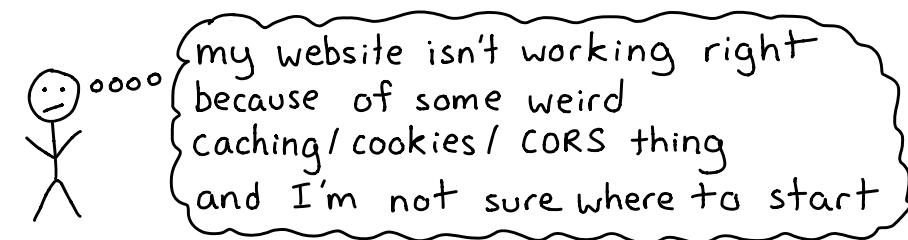


Like this?
more zines at
wizardzines.com

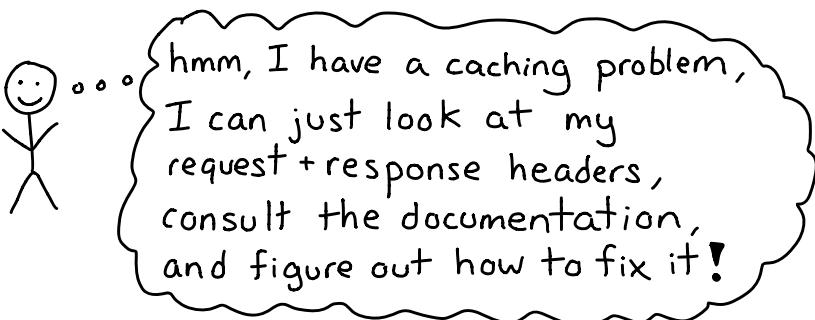
about this zine

Your browser uses HTTP every time it visits a website. Like a lot of the tech that runs the internet, understanding HTTP isn't that hard!

This zine's goal is to take you from:



to



credits

Cover art: Vladimir Kašiković
Editing: Dolly Lanuza, Kamal Marhubi
special thanks to Marco Rogers for suggesting the idea of a HTTP zine

how to learn more

♥ Mozilla Developer Network

<https://developer.mozilla.org>

MDN is a fantastic wiki maintained by Mozilla. It has tutorials and reference documentation for HTML, CSS, HTTP, Javascript. It's the best place to start for reference documentation.

♥ OWASP

<https://cheatsheetseries.owasp.org>

OWASP is an organization that publishes security best practices. If you have a question about web security, they've probably published a cheat sheet or guide to help you.

♥ httpstatuses.com

Nice little site that explains all the HTTP status codes.

♥ RFCs

<https://tools.ietf.org/html/rfcXXXX>

put RFC number here

RFCs are numbered documents (like "RFC 2631"). Every Internet protocol (like TLS or HTTP) has an RFC. These are where you go to find the Official Final Answers to technical questions you have about any internet standard. The HTTP standard is mostly documented in 6 RFCs numbered 7230 to 7235.



is the Host header
actually required?



Yes, section 5.4
of RFC 7230
says so!

the final answer
Don't be scared of using an RFC if you want to know for sure!

Table of Contents

What's HTTP?	4
How URLs work	5
What's a header?	6
Request methods	7
Anatomy of an HTTP request	8-9
Request methods (GET! POST!)	8-9
Using HTTP APIs	10
Request headers	11
Responses:	12
Anatomy of an HTTP response	13
Status codes (200! 404!)	14
How cookies work	15
Content delivery networks & caching	16-17
HTTP/2	18
Redirections	19
HTTP/2	20-21
Same origin policy & CORS	22-24
Security headers	25
Exercises & how to learn more	26-27

Making HTTP requests with curl to real internet websites and trying different headers is my favourite way to play around with HTTP & learn.

-i shows the response headers
-H adds a request header
-I only shows the response headers by sending a HEAD request

* curl tips

curl -i https://examplecat.com/cat.txt -H "Range: bytes=8-17"

Try the Range header:

curl -i https://examplecat.com "Accept-Language: es-ES"

Get a webpage in Spanish:

curl -i http://twitter.com -H "Accept-Encoding: gzip" -- output -

Request (and print out!) a compressed response:

curl -i https://examplecat.com "Accept-Language: es-ES"

Get redirected to another URL:

curl -i http://location header) (hint: look at the

Guess what content delivery network GitHub is using:

curl -I https://github.githubbasestats.com (hint: it's in a header starting with x-)

Find out when example.com was last updated

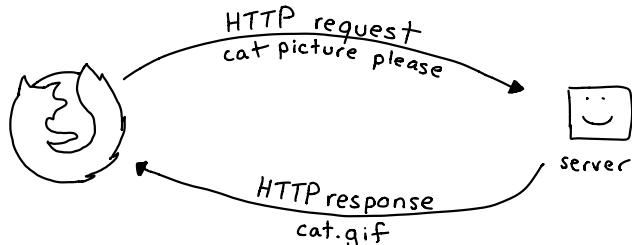
curl -I example.com (hint: Last-Modified)

Get a 404 not found

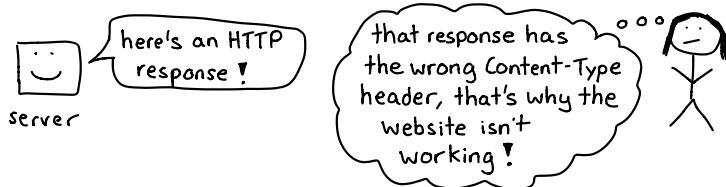
curl -i example.com/bananas (Get a 404 not found)

what's HTTP?

HTTP is the protocol (Hypertext Transfer Protocol) that's used when you visit any website in your browser.



The exciting thing about HTTP is that even though it's used for literally every website, HTTP requests and responses are easy to look at and understand:



Example of what an HTTP request and response might look like:

request	response
request line headers { Host: examplecat.com User-Agent: curl Accept: */*	status { HTTP/1.1 200 OK headers { Cache-Control: max-age=604800 Content-Type: text/html Etag: "1541025663+ident" Server: ECS (nyb/1D0B) Vary: Accept-Encoding X-Cache: HIT Content-Length: 1270 body { <!doctype html> <title>Example Cat</title> ...

All that text is a lot to understand, so let's get started learning what all of it means!

security headers

These are headers your server can set. They ask the browser to protect your users' data against attackers in different ways:

Content-Security-Policy often called CSP

Only allow CSS / Javascript from certain domains you choose to run on your website. Helps protect against cross-site-scripting (aka XSS) attacks.

Referrer-Policy

Control how much information is sent to other sites in the Referer header. Example: Referrer-Policy: no-referrer.
spelling is inconsistent with Referer header !!

Strict-Transport-Security often called HSTS

Require HTTPS. If you set this the client (browser) will never request a plain HTTP version of your site again. Be careful! You can't take it back!

Expect-CT

Certificate Transparency (CT) is a system that can help find malicious SSL certificates issued for your site. This header gives the browser a URL to use to report bad certificates to you.

X-XSS-Protection

Another way to protect against XSS attacks. Not supported by all browsers, Content-Security-Policy is more powerful.

5

scheme	<code>https://examplecat.com:443/cats?color=light%20gray#banana</code>	domain	port	path	query string	fragment id
protocol	Protocol to use for the request. Encrypted (https),	where to send the request.	for HTTP(s) requests,	the Host header gets set to this (Host: example.com)	examplecat.com	domain
insecure	insecure (http), or something else entirely (ftp).	Defaults to 80 for HTTP and 443 for HTTPS.	Path to ask the server for. The path and the	a different version of a page ("I want a light	/cats/light	Path
query parameters	Query parameters are usually used to ask for	query parameters are combined in the request,	like: GET /cats?color=light%20gray HTTP/1.1	gray cat!!"). Example:	color=light gray	query parameters
URL encoding	URLs aren't allowed to have certain special	characters like spaces, @, etc. So to put them in a	URL you need to percent encode them as	% + hex representation of ASCII value.	name = value separated by &	URL
URLEncoder	URLs aren't allowed to have certain special	characters like spaces, @, etc. So to put them in a	URL you need to percent encode them as	%20	URLEncoder	URLEncoder
JavaScript on the page.	This isn't sent to the server at all. It's used either	to jump to an HTML tag () or by	#banana	Javaascrip	fragme	JavaScript

How URLs work

Cross-Origin Resource Sharing (CORS)

CORS

cross-origin resource sharing

Cross-origin requests are not allowed by default:
(because of the same origin policy!)

POST requests to api.clothes.com?
is a different origin from clothes.com
NOPE. api.clothes.com?

clothes.com

script from clothes.com

same origin flowchart

If you run `api.clothes.com`, you can allow `clothes.com` to make requests to it using the `Access-Control-Allow-Origin` header.

This diagram illustrates the flow of an OPTIONS preflight request between a client and a server.

Client (Left):

- Says: "That's cross-origin. I'm going to need to ask `api.clothes.com` if this request is allowed."
- Asks: `POST /buy-things` (Host: `api.clothes.com`)

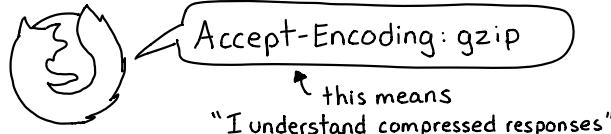
Server (Right):

- Answers: "Hey, what requests are allowed?" (Host: `api.clothes.com`)
- Replies: `OPTIONS /buy-things` (Host: `api.clothes.com`)
- Shows: "I'm allowed to make requests to `api.clothes.com` if this request is allowed."
- Shows: "Cool, the request is allowed!" (Host: `api.clothes.com`)
- Shows: `204 No Content` (Host: `api.clothes.com`)
- Shows: "Access-Control-Allow-Origin: clothes.com" (Host: `api.clothes.com`)
- Shows: "200 OK" (Host: `api.clothes.com`)
- Shows: "POST /buy-things" (Host: `api.clothes.com`)
- Shows: "Referer: api.clothes.com/checkout" (Host: `api.clothes.com`)
- Shows: "POST /buy-things-bought" (Host: `api.clothes.com`)
- Shows: "[{"things-bought": true}]"

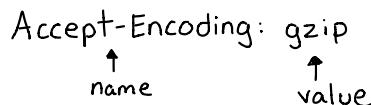
requests that send JSON need a preflight

what's a header?

Every HTTP request and response has headers. Headers are a way for the browser or server to send extra information!



Headers have a name and a value.



Header names aren't case sensitive:

totally valid ↗ aCcEPT-eNcOdInG : gzip

There are a few different kinds of headers:

Describe the body:

Content-Type: image/png Content-Encoding: gzip
Content-Length: 12345 Content-Language: es-ES

Ask for a specific kind of response:

Accept: image/png
Range: bytes=1-10

Accept-Encoding: gzip
Accept-Language: es-ES

Every Accept-header has a corresponding Content- header

Manage caches:

ETag: "abc123"
If-None-Match: "abc123"
Vary: Accept-Encoding
If-Modified-Since: 3 Aug 2019 13:00:00 GMT
Last-Modified: 3 Feb 2018 11:00:00 GMT
Expires: 27 Sep 2019 13:07:49 GMT
Cache-Control: public, max-age=300

Say where the request comes from:

User-Agent: curl

Referer: https://examplecat.com

Cookies:

Set-Cookie: name=julia; HttpOnly (server → client)
Cookie: name=julia (client → server)

6

and more!

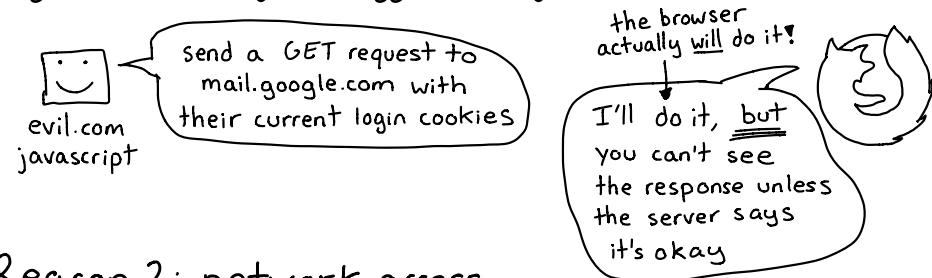
why the same origin policy matters

Browsers work hard to make sure that evil.com can't make requests to other-website.com. But evil.com can request other-website.com from its own server, what's the big deal?

2 reasons it's important to restrict Javascript on websites from making arbitrary requests from your browser:

Reason 1: cookies

Browsers often send your cookies with HTTP requests. You don't want evil.com to be able to make requests using your login cookies. They'd be logged in as you!



Reason 2: network access

You might be on a private network (for example your company's corporate network) that evil.com doesn't have access to, but your computer does.



23

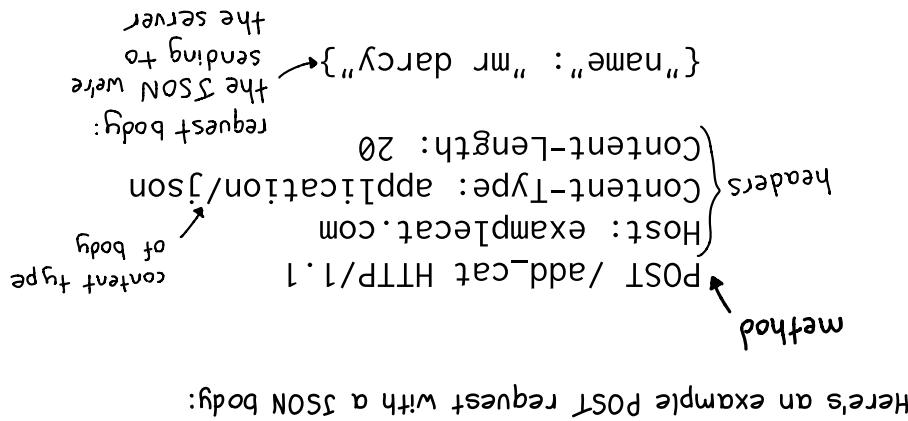
* Anatomy of an HTTP request *

HTTP requests always have:

- a domain (like example.com)
- a resource (like /cat.jpg)
- a method (GET, POST, or something else)
- headers (extra information for the server)
 - ↳ a resource being requested
 - ↳ usually GET or POST
 - ↳ HTTP version
 - ↳ domain being requested
 - ↳ URL in your browser. It doesn't have a body.
- This is an HTTP 1.1 request for example.cat.com/cat.jpg.
- It's a GET request, which is what happens when you type a URL in your browser. It doesn't have a body.
- This is an optional request body. GET requests usually don't have a body, and POST requests usually do.

```

graph TD
    A[HTTP requests always have:] --> B["→ a domain (like example.com)"]
    A --> C["→ a resource (like /cat.jpg)"]
    A --> D["→ a method (GET, POST, or something else)"]
    A --> E["→ headers (extra information for the server)"]
    E --> F["↳ a resource being requested"]
    E --> G["↳ usually GET or POST"]
    E --> H["↳ HTTP version"]
    E --> I["↳ domain being requested"]
    E --> J["↳ URL in your browser. It doesn't have a body."]
    E --> K["→ This is an optional request body"]
    E --> L["→ It's a GET request, which is what happens when you type a URL in your browser. It doesn't have a body."]
    E --> M["→ This is an HTTP 1.1 request for example.cat.com/cat.jpg."]
    E --> N["→ It's a GET request, which is what happens when you type a URL in your browser. It doesn't have a body."]
    E --> O["→ This is an optional request body"]
    E --> P["→ GET requests usually don't have a body,"]
  
```



```

graph TD
    A["An origin is the protocol + domain including subdomains + port"] --> B["The same origin policy is one way browsers protect you from malicious JavaScript code. Here's basically how it works:"]
    B --> C["example: https://babby.example.com:443"]
    C --> D["please make a request to this URL?"]
    D --> E["let me check my charset!"]
    E --> F["does the origin match exactly?"]
    F --> G["allows CORS time?"]
    G --> H["is this request type allowed"]
    H --> I["from a different origin?"]
    I --> J["CSS, <img src>, and a few other things are OK"]
    J --> K["is this request type allowed"]
    K --> L["allowed"]
    L --> M["is no, not yes?"]
    M --> N["notice the default"]
    N --> O["is no, not yes?"]
    O --> P["do an options preflight request"]
    P --> Q["do the response's CORS headers allow this?"]
    Q --> R["do the response's CORS headers allow this?"]
    R --> S["DENIED"]
    R --> T["HEADERS ALLOWED"]
  
```

The same origin policy is one way browsers protect you from malicious JavaScript code. Here's basically how it works:

example: `https://babby.example.com:443`

`please make a request to this URL?`

`let me check my charset!`

`does the origin match exactly?`

`allows CORS time?`

`is this request type allowed`

`from a different origin?`

`CSS, , and a few other things are OK`

`is no, not yes?`

`notice the default`

`is no, not yes?`

`do an options preflight request`

`do the response's CORS headers allow this?`

`do the response's CORS headers allow this?`

`HEADERS DENIED`

`HEADERS ALLOWED`

request methods

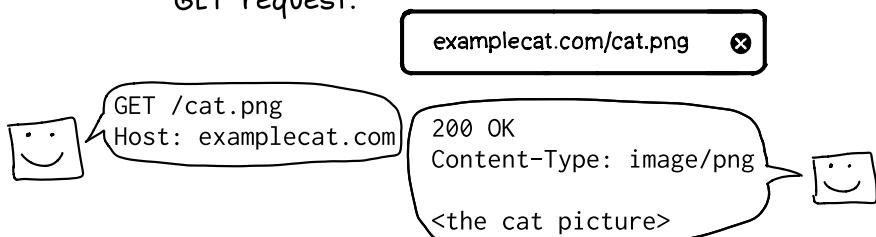
Every HTTP request has a method. It's the first thing in the first line:

 this means it's a *GET* request

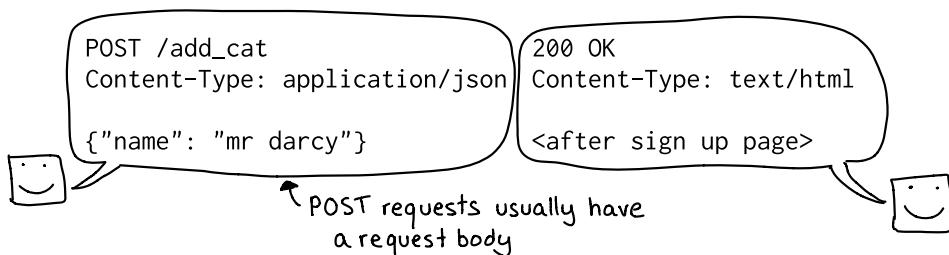
GET /cat.png HTTP/1.1

There are 9 methods in the HTTP standard. 80% of the time you'll only use 2 (GET and POST).

GET When you type an URL into your browser, that's a GET request.

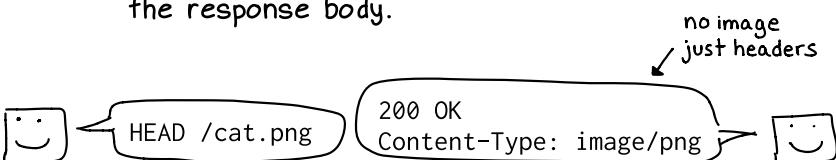


POST When you hit submit on a form, that's (usually) a POST request.



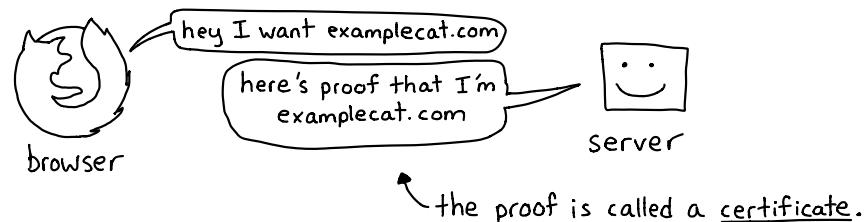
The big difference between GET and POST is that GETs are never supposed to change anything on the server.

HEAD Returns the same result as GET, but without the response body.



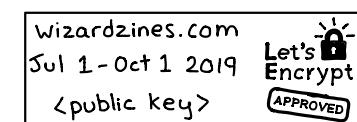
certificates

To establish an HTTPS connection to examplecat.com, the client needs proof that the server actually is examplecat.com.



A TLS certificate has:

- a set of domains it's valid for (eg examplecat.com)
- a start and end date (example: july 1 2019 to oct 1 2019)
- a secret private key which only the server has this is the only secret part, the rest is public
- a public key to use when encrypting
- a cryptographic signature from someone trusted

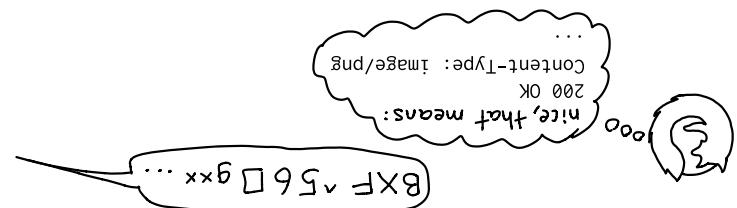


The trusted entity that signs the certificate is called a ★ Certificate Authority ★ (CA) and they're responsible for only signing certificates for a domain for that domain's owner.

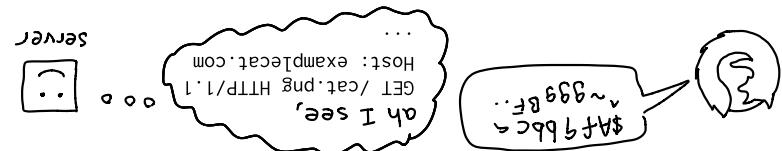


When your browser connects to examplecat.com, it validates the certificates using a list of trusted CAs installed on your computer. These CAs are called "root certificate authorities".

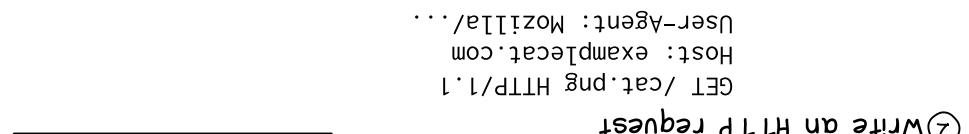




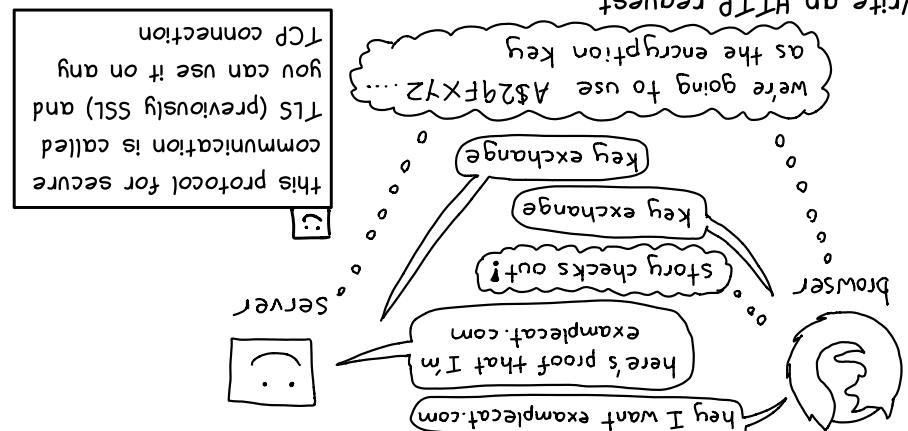
④ Receive encrypted HTTP response



③ Encrypt the HTTP request with AES & send it to examplecat.com



② Write an HTTP request



Simplified version of how picking the encryption key works:

server will use the same key to encrypt/decrypt content.

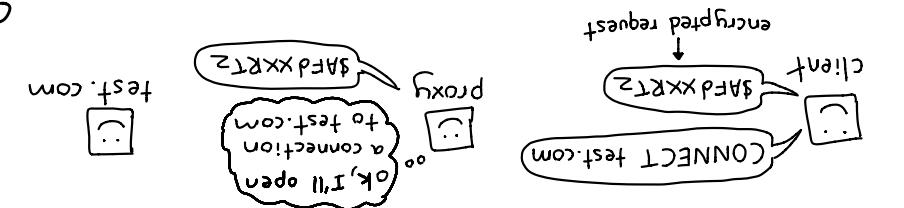
use for this connection to examplecat.com. The browser and server will use the same key to encrypt/decrypt content.

① Negotiate an encryption key (AES symmetric key) to

<https://examplecat.com/cat.png>:

Here's what your browser does when it asks for

HTTPS: HTTP + Secure ⚡



will use this protocol to proxy your requests.

variable to a proxy server, many HTTP libraries

If you set the HTTPS_PROXY environment

proxy to open a connection.

a request to a server directly, it asks for a

Different from the others: instead of making

you probably don't need to know about it.

I've never seen a server that supports this,

resource ("just change this file").

Used in some APIs for partial updates to a

GET /cat/1234 later.

or update resources. PUT /cat/1234 lets you

Used in some APIs (like the S3 API) to create

GET /cat/1234 later.

TRACE

PATCH

PUT

DELETE

Used in many APIs (like the Stripe API) to

delete resources.

It also tells you which methods are available.

The CORS page has more about that.

OPTIONS is mostly used for CORS requests.

OPTIONS

request headers

These are the most important request headers:

Host

The domain.
The only required header.

Host: examplecat.com User-Agent: curl 7.0.2 Referer: https://examplecat.com

User-Agent

Name + version of your browser and OS

User-Agent: curl 7.0.2 Referer: https://examplecat.com

Referer

website that linked or included the resource

Referer: https://examplecat.com
↑ yes, it's misspelled!

Authorization

eg a password or API token
base64 encoded user:password

Authorization: Basic YXZ

Cookie

Send cookies the server sent earlier
keeps you logged in.

Cookie: user=b0rk

Range

lets you continue downloads ("get bytes 100-200")
Range: bytes=100-200

Cache-Control

"max-age = 60"
means cached responses must be less than 60 seconds old

If-Modified-Since: Wed, 21 Oct...

Accept

MIME type you want the response to be

Accept: image/png

If-Modified-Since

only send if resource was modified after this time

only send if the ETag doesn't match those listed

If-None-Match: "e7ddac"

Accept-Encoding

set this to "gzip" and you'll probably get a compressed response

Accept-Encoding: gzip

Accept-Language

set this to "fr-CA" and you might get a response in French

Accept-Language: fr-CA

Content-Type

MIME type of request body, e.g. "application/json"

Content-Encoding

will be "gzip" if the request body is gzipped

Connection

"close" or "keep-alive". Whether to keep the TCP connection open.

HTTP/2

HTTP/2 is a new version of HTTP.

Here's what you need to know:

★ A lot isn't changing

All the methods, status codes, request/response bodies, and headers mean exactly the same thing in HTTP/2.

before (HTTP/1.1)

method: GET
path: /cat.gif
headers:
- Host: examplecat.com
- User-Agent: curl

after (HTTP/2)

method: GET
path: /cat.gif
headers:
- User-Agent: curl
one change:
Host header
=> authority
authority: examplecat.com

★ HTTP/2 is faster

Even though the data sent is the same, the way HTTP/2 sends it is different. The main differences are:

- It's a binary format (it's harder to tcpdump traffic and debug)
- Headers are compressed
- Multiple requests can be sent on the same connection at a time

before (HTTP/1.1)

→ request 1
→ request 2
→ response 1 ←
→ response 2 ←

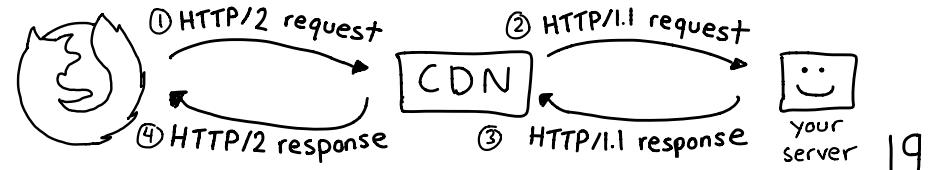
after (HTTP/2)

→ request 1
→ request 2
out of order { response 2 ←
response 1 ← }
one TCP connection

All these changes together mean that HTTP/2 requests often take less time than the same HTTP/1.1 requests.

★ Sometimes you can switch to it easily

A lot of software (CDNs, nginx) let clients connect with HTTP/2 even if your server still only supports HTTP/1.1.



Using HTTP APIs

• Lots of services (Twitter! Twilio! Google!) let you use them by sending them HTTP requests. If an HTTP API doesn't come with a client library, don't be scared! You can just make the HTTP requests yourself. Here's what you need to remember:

Set the right {Content-Type} header
Often you'll be sending a POST request with a body, and that means you need a Content-Type header that matches the body

- The 2 main options are:
 - * application/json → JSON!
 - * application/x-www-form-urlencoded → same as what a HTML form does

A common error is to try to send POST data as one content type (like JSON) when it's actually another (like application/x-www-form-urlencoded)

Most HTTP APIs require a secret API key so they know who you are.

```
curl https://api.twilio.com/2010-04-01/Accounts/ACCOUNT_ID/Messages.json  
-H "Content-Type: application/json"  
-u ACCOUNT_ID:AUTH_TOKEN  
-d {  
    "from": "+15141234567",  
    "to": "+15141234567",  
    "body": "a text message"  
}
```

redirects

Sometimes you type a URL into your browser:

but end up at a slightly different URL:
ooh, where did the cat come from?

Here's what's going on behind the scenes:

The diagram illustrates a 301 Moved Permanently response. It shows a browser sending a GET request to a server. The server returns a 301 Moved Permanently response with a Location header pointing to the new URL. The browser then sends a new GET request to the new URL, which the server handles normally.

```
graph LR; Browser[...] -- "GET /cat.png HTTP/1.1" --> Server[...]; Server -- "HTTP/1.1 301 Moved Permanently<br/>Location: /cat.png" --> Browser; Browser -- "GET /cat.png HTTP/1.1" --> Server[...]; Server -- "HTTP/1.1 200 OK<br/>Host: example.com<br/>200 OK<br/><rest of website here>" --> Browser;
```

The `Location` header tells the browser what new URL to use. The new URL doesn't have to be on the same domain:

```
example.com/panda can redirect to padadas.com.
```

Setting up redirects is a great thing to do if you move your site

anatomy of an HTTP response

HTTP responses have:

- a status code (200 OK! 404 not found!)
- headers
- a body (HTML, an image, JSON, etc)

Here's the HTTP response from examplecat.com/cat.txt:

```

HTTP/1.1 200 OK status
Accept-Ranges: bytes
Cache-Control: public, max-age=0
Content-Length: 33
Content-Type: text/plain; charset=UTF-8
Date: Mon, 09 Sep 2019 01:57:35 GMT
Etag: "ac5affa59f554a1440043537ae973790-ssl"
Strict-Transport-Security: max-age=31536000
Age: 0
Server: Netlify

\   /
) ( ' ) cat! ^
( / )
\(_)
  
```

} status code

} headers

} body

There are a few kinds of response headers:

when the resource was sent/modified:

Date: Mon, 09 Sep 2019 01:57:35 GMT
Last-Modified: 3 Feb 2017 13:00:00 GMT

about the response body:

Content-Language: en-US Content-Type: text/plain; charset=UTF-8
Content-Length: 33 Content-Encoding: gzip

caching:

ETag: "ac5affa..." Age: 255
Vary: Accept-Encoding Cache-Control: public, max-age=0
security: (see page 25)

X-Frame-Options: DENY Strict-Transport-Security: max-age=31536000
X-XSS-Protection: 1 Content-Security-Policy: default-src https:

and more:

Connection: keep-alive Accept-Ranges: bytes
Via: nginx
Set-Cookie: cat=darcy; HttpOnly; expires=27-Feb-2020 13:18:57 GMT;

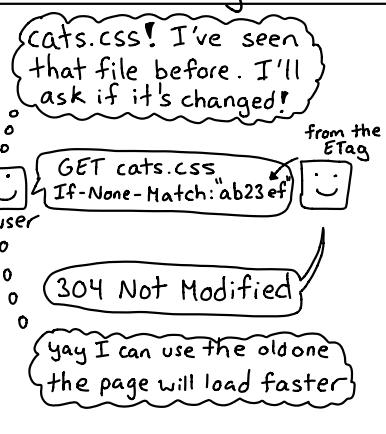
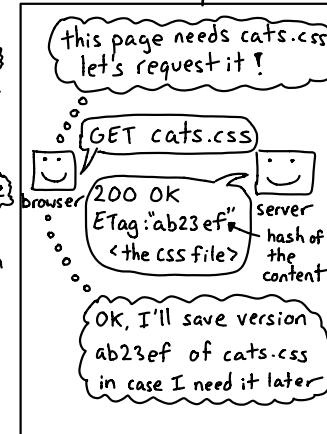
caching headers

These 3 headers let the browser avoid downloading an unchanged file a second time.

ETag
response header
and

If-None-Match
request header

If-Modified-Since
is similar to
If-None-Match
but with
Last-Modified
instead of ETag



Vary
response header

Sometimes the same URL can have multiple versions (spanish, compressed or not, etc). Caches categorize the versions by request header like this:

Accept-Language	Accept-Encoding	content
en-US	-	hello
es-ES	-	hola
en-US	gzip	f\$xx99aef^.. (compressed gibberish)

The Vary header tells the cache which request headers should be the columns of this table.

Cache-Control

request AND response header

Used by both clients and servers to control caching behaviour. For example:

Cache-Control: max-age=999999999999 from the server asks the CDN or browser to cache the thing for a long time.

is supported for this resource
Whether Range requests header

Accept - Ranges

Whether body is compressed
Content-Encoding: gzip

Content - Encoding

Length of body in bytes
Content-Length: 33

Content - Length

allow cross-origin requests.
Cross-Origin Resource Sharing (CORS) headers. These

Access - Control - *

Keep-alive
TCP connection open
Whether to keep the

Connection

That response will
vary based on
request headers

Vary

When content was
last modified
(not always accurate)

Last - Modified

response headers

URL to redirect to
Location: /cat.png

Location

Language of body
Content-Language: en-US

Content - Language

MIME type of body
Content-Type: text/plain

Content - Type

Sets a cookie.
Set-Cookie: name=value; HttpOnly

Set - Cookie

Proxy servers
added by
via: negotia
and should be re-requested
The response is stable
after this time.

Expires

Cache-Control
max-age=300
Cache-control
max-age=300

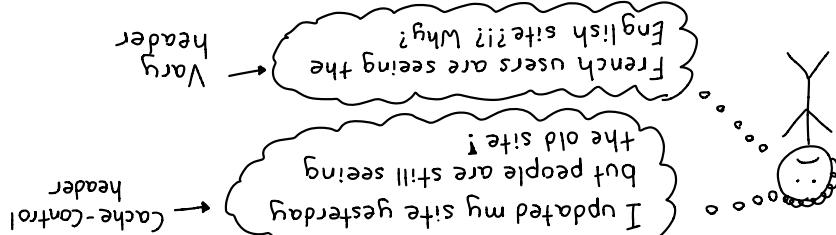
Etag: "ac5afffa.."
Response body
version of
responsible
set-timing
various caching

Etag

When response
was sent
Date: Mon, 09 Sep 2019..
Age: 355
seconds response
how many

Age

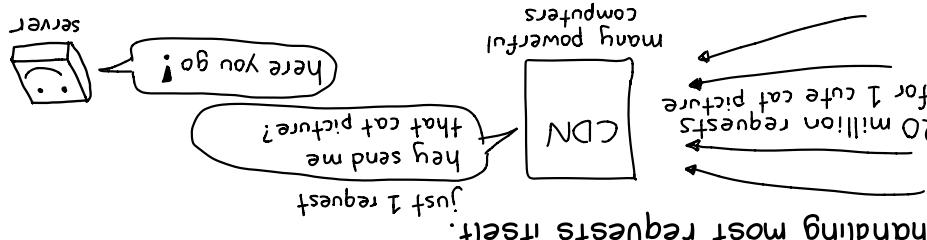
Next, we'll explain the HTTP headers your CDN or browser uses to decide how to do caching.



This is great but caching can cause problems too!

Today, there are many free or cheap CDN services available, which means if your site gets popular suddenly you can easily keep it running!

Up until now, you can easily keep it running!



Handling most requests itself.

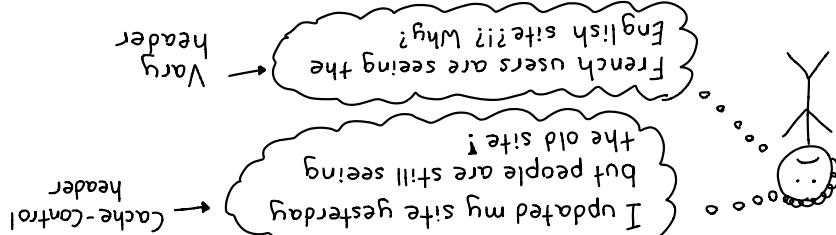
A CDN (content delivery network) can make your site faster and save you money by caching your site and



In 2004, if your website suddenly got popular, often the webserver wouldn't be able to handle all the requests.

content delivery networks

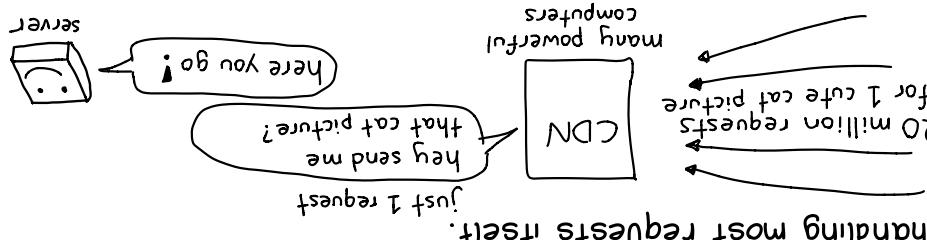
Next, we'll explain the HTTP headers your CDN or browser uses to decide how to do caching.



This is great but caching can cause problems too!

Today, there are many free or cheap CDN services available, which means if your site gets popular suddenly you can easily keep it running!

Up until now, you can easily keep it running!



Handling most requests itself.

A CDN (content delivery network) can make your site faster and save you money by caching your site and



In 2004, if your website suddenly got popular, often the webserver wouldn't be able to handle all the requests.

HTTP status codes

Every HTTP response has a ★status code★.



There are 50ish status codes but these are the most common ones in real life:

200 OK

301 Moved Permanently

302 Found

temporary redirect

304 Not Modified

the client already has the latest version, "redirect" to that

400 Bad Request

403 Forbidden

API key/OAuth/something needed

404 Not Found

we all know this one :)

429 Too Many Requests

you're being rate limited

500 Internal Server Error

the server code has an error

503 Service Unavailable

could mean nginx (or whatever proxy)

couldn't connect to the server

504 Gateway Timeout

the server was too slow to respond

} 2xx s mean
★ success ★

} 3xx s aren't errors, just redirects to somewhere else

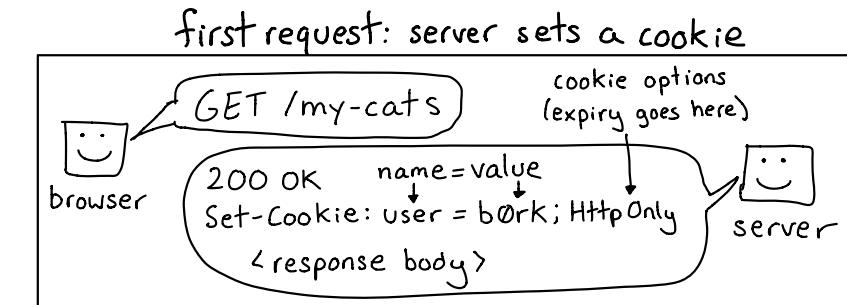
} 4xx errors are generally the client's fault:
it made some kind of invalid request

} 5xx errors generally mean something's wrong with the server.

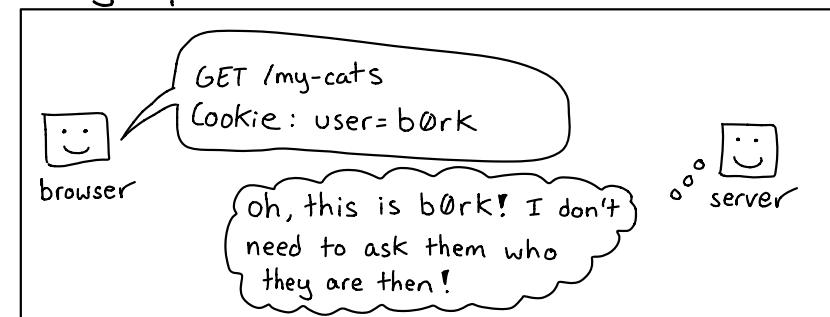
how cookies work

Cookies are a way for a server to store a little bit of information in your browser.

They're set with the Set-Cookie response header, like this:



Every request after: browser sends the cookie back



Cookies are used by many websites to keep you logged in. Instead of `user=b0rk` they'll set a cookie like `sessionid=long-incomprehensible-id`. This is important because if they just set a simple cookie like `user=b0rk`, anyone could pretend to be b0rk by setting that cookie!

Designing a secure login system with cookies is quite difficult — to learn more about it, google "OWASP Session Management Cheat Sheet".