

love this?

find more awesome zines at

→ jvns.ca/zines ←

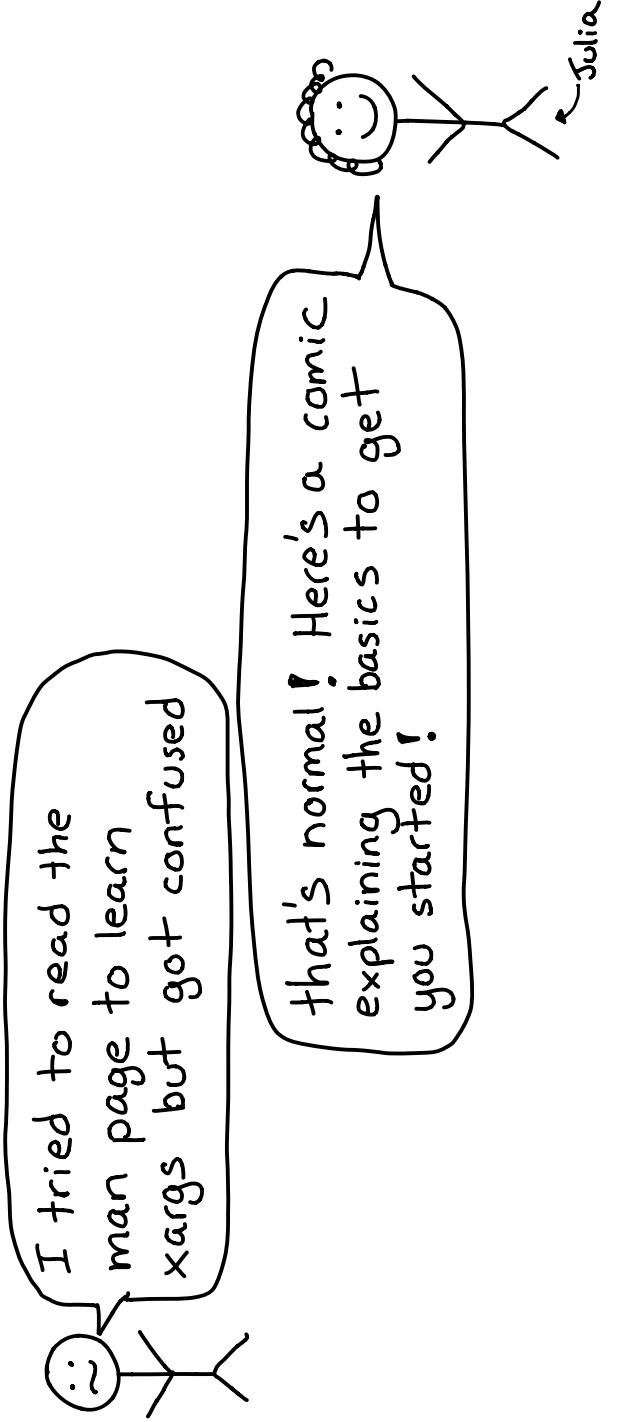
CC-BY-NC-SA computer wizard industries 2018

By Julia Evans

Command Line Size Sizing Rite



This zine explains some of the most useful Unix command line tools in 1 page each.



Even if you've used the command before, I might have a new trick or two for you ❤

*

more useful tools

- make
- jq
- nohup
- disown
- cut/paste
- sponge
- xxd
- hexdump
- objdump
- strings
- screen
- tmux
- date
- entr
- seq
- join
- parallel:
 - GNU parallel
 - pigg / pixz
 - sort --parallel
- diff -U
- vipe
- imagemagick
- fish
- ranger
- chronic

*

Isof

Isof

what Isof tells you

for each open file:

→ pid
→ file type (regular? directory?
FIFO? socket?)

→ file descriptor (FD column)
→ user
→ filename / socket address

Isof

stands for list open files



I can tell you!

-i

list open network sockets
(sockets are files!)

examples:

- i -n -P ← "-n & -P mean "don't resolve"
- i :8080 host names / ports" (also -Pnij)
- i TCP
- i -s TCP:LISTEN

find deleted files

\$ Isof | grep deleted

will show you deleted files!
You can recover open deleted
files from
/proc/<pid>/fd/<fd>

processes that opened the file

netstat

another way to list open
sockets on Linux is:
netstat -tunapl
↑
tuna, please!

On Mac netstat has
different args.

-P PID

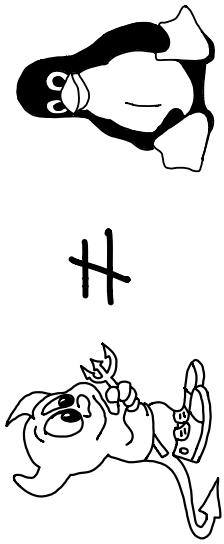
list the files PID has
open

Isof /some /dir

list just the open files
in /some /dir

Table of contents

BSD#GNU.....4	bash tricks.....10-11	misc commands...17
grep.....5	disk usage.....12	head & tail.....18
find.....6	tar.....13	less.....19
xargs.....7	ps.....14	kill.....20
awk.....8	top.....15	cat.....21
sed.....9	sort & uniq.....16	Isof.....22

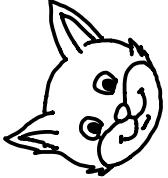


For almost all these tools, there are at least 2 versions:

① The BSD version (on BSDs & Mac OS)

② The GNU version (on Linux)

All of the examples in this zine were tested on Linux. Some things (like sed -i) are different on Mac. Be careful when writing cross-platform scripts! You can install the GNU versions on Mac with 'brew install coreutils'.



cat & friends

21

cat concatenates files

```
$ cat myfile.txt
prints contents of myfile.txt
$ cat *.txt
prints all .txt files put together!
```

you can use cat as an EXTREMELY BASIC text editor:

- ① Run \$ cat > file.txt
- ② type the contents (don't make mistakes !)
- ③ press ctrl-d to finish

cat -n
prints out the file with line numbers!

- 1 Once upon a midnight...
- 2 Over many a quaint...
- 3 While I nodded, nearly

zcat

cats a gzipped file!
Actually just a 1-line shell script that runs 'gzip -cd', but easier to remember.

tee

'tee file.txt' will write its stdin to both stdout and file.txt
→ stdout → a.txt
stdin → tee a.txt → a.txt

how to redirect to a file owned by root
\$ sudo echo "hi">>x.txt

↑
this will open x.txt as your user, not as root, so it fails!
\$ echo "hi" | sudo tee -a x.txt
will open x.txt as root ↴

Kill

kill doesn't just kill programs



kill - SIGNAL PID

name or number

killall - SIGNAL NAME

signals all processes called NAME for example

\$ killall firefox

useful flags:

- w wait for all signaled processes to die
- i ask before signalling

which signal kill sends	<u>name</u>	<u>num</u>
kill	=> SIGTERM	15
kill -q	=>	SIGKILL 9
kill -KILL	=>	SIGKILL 9
kill -HUP	=>	SIGHUP 1
kill -STOP	=>	SIGSTOP 19
POLL	30	PUR 31
TRAP	5	SIGTRAP 4
ABRT	6	SIGABRT 7
ALRM	14	SIGALRM 15
CONT	18	SIGCONT 19
CONT	19	SIGSTOP 20
TSTP	20	SIGTSTP 21
TTIN	22	SIGTTIN 23
TTOUT	23	SIGTTOUT 24
URG	25	SIGURG 25
VTIME	26	SIGVTIME 27
VTIME	27	SIGPROF 28
WINCH	28	SIGWINCH 29

pgrep

prints PIDs of matching running programs

pgrep fire matches firefox firebird

NOR bash firefox.sh

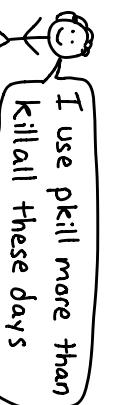
To search the whole command line (eg bash firefox.sh) use {pgrep -f}

Kill -R	lists all signals
HUP	SIGINT 2
TRAP	SIGQUIT 3
ABRT	SIGILL 4
USR1	SIGPOLL 5
SEGV	SIGPOLL 6
USR2	SIGPOLL 7
PIPE	SIGPOLL 8
ALRM	SIGPOLL 9
TERM	SIGPOLL 10
STKFLT	SIGPOLL 11
CONT	SIGPOLL 12
STOP	SIGPOLL 13
TSTP	SIGPOLL 14
TTIN	SIGPOLL 15
TTOUT	SIGPOLL 16
URG	SIGPOLL 17
VTIMER	SIGPOLL 18
PROF	SIGPOLL 19
WINCH	SIGPOLL 20
POLL	SIGPOLL 21
PUR	SIGPOLL 22
SYS	SIGPOLL 23

pkill

same as pgrep, but signals PIDs found. ex:

pkill -f firefox



grep

5

Recursive! Search all the files in a directory.

-o only print the matching part of the line (not the whole line)

-a search binaries: treat binary data like it's text instead of ignoring it!

Grep alternatives **ack** **ag** **ripgrep** (better for searching code!)

-A show context for your search.
\$ grep -A 3 foo will show 3 lines of context after a match

-B
 -C
\$ grep -B 3 foo
will show 3 lines of context before a match

-i case insensitive

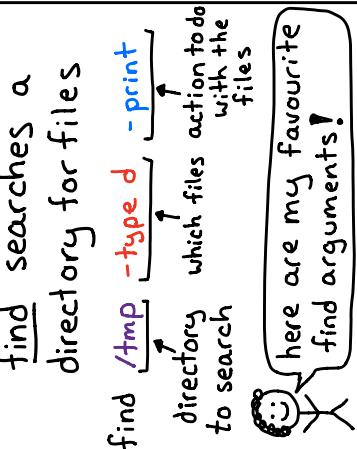
-l only show the filenames of the files that matched

-F aka **grep** don't treat the match string as a regex eg \$ grep -F ...

find

6

-name/-iname case insensitive the filename! eg <code>-name '* .txt'</code>	-type [TYPE] f: regular file l: symlink d: directory + more!
-path / -ipath search the full path! <code>-path '/home/*/*.go'</code>	-maxdepth NUM only descend NUM levels when searching a directory
-print0 print null-separated filenames Use with xargs -0!	-locate The locate command searches a database of every file on your system. good: faster than find bad: can get out of date \$sudo updatedb updates the database
-size 0 find empty files! Useful to find files you created by accident	-delete action: delete all files found
-exec COMMAND action: run COMMAND on every file found	



19

less

less is a pager that means it lets you view (not edit) text files or piped in text man uses your pager (usually less) to display man pages	q uit V lowercase edit file in your \$EDITOR	F press F to keep reading from the file as it's updated (like tail -f) press Ctrl+C to stop reading updates
	+ + runs a command when less starts	
	less +F : follow updates	
	less +G : start at end of file	
	less +20% : start 20% into file	
	less +/foo : search for 'foo' right away	

head & tail

18

head

shows you the first 10 lines of a file

if you pipe a program's output to head, the program will stop after printing 10 lines (it gets sent SIGPIPE)

-c NUM

show the first/last NUM bytes of the file

head -c 1K
will show the first 1024 bytes

tail

tail shows the last 10 lines!

tail -f FILE will follow:
print any new lines added to the end of FILE. Super useful for log files!

tail --retry

keep trying to open file if it's inaccessible

tail --pid PID
stop when process PID stops running (with -f)

-n NUM

-n NUM (either head or tail)
will change the # lines shown

head -n -NUM } show all tail -n +NUM } last / first NUM lines

Xargs

7

xargs takes whitespace separated strings from stdin and converts them into command-line arguments

```
$ echo "/home /tmp" | xargs ls  
will run  
ls /home /tmp
```

this is useful when you want to run the same command on a list of files!

```
→ delete (xargs rm)  
→ combine (xargs cat)  
→ search (xargs grep)  
→ replace (xargs sed)
```

how to replace "foo" with "bar" in all .txt files:

```
find . -name '*.txt' |  
xargs sed -i s/foo/bar/g
```

how to lint every Python file in your Git repo:

```
git ls-files | grep .py |  
xargs pep8
```

if there are spaces in your filenames "my day.txt"
xargs will think it's 2 files "my" and "day.txt"

fix it like this:

```
find . -print0 |  
xargs -0 COMMAND
```

more useful xargs options

-n 1 makes xargs run a separate process max-args for each input.
-P is the max number of parallel processes xargs will start

awk

8

awk is a tiny programming language for manipulating columns of data

I only know how to do 2 things with awk but it's still useful!

basic awk program structure

```
BEGIN{ ... } {action}
CONDITION {action}
CONDITION {action}
END{ ... } {action on lines matching CONDITION}
```

extract a column of text with awk

```
awk -F, ' {print $5}'
```

↑
column separator quotes! 5th column
this is 99% of what I do with awk

so many unix commands print columns of text (ps! ls!) so being able to get the column you want with awk is GREAT

awk program example: sum the numbers in the 3rd column

```
{action
{s += $3}; END{print s}}
at the end, print the sum!
```

awk program example: print every line over 80 characters

```
[length($0) > 80]
{condition
(there's an implicit {print} as the action)}
```

misc commands

17

rlwrap

adds history & ctrl support to REPLs that don't already have them (rl stands for readline)

```
$ rlwrap python
```

watch

rerun a command every 2 seconds

file

figures out what kind of file (png? pdf?) a file is

+s

add a timestamp in front of every input line

comm

find lines 2 sorted files have in common

pv

"pipe viewer", gives you stats on data going through a pipe

cal

a tiny calendar ☺

diff

diff 2 files. Run with '-u 8' for context.

column

format input into columns

xsel / xclip copy/paste from system clipboard.
(pbcopy / pbpaste on Mac)

Sort & uniq

sort sorts its input

\$ sort names.txt
the default sort is alphabetical.

uniq removes duplicates

a
b
b
a
a
c
c
notice there
are still 2
'a's! uniq
only uniqueness
adjacent
matching lines

sort -n
numeric sort
'sort' order 'sort-n' order
12 12
15000 48
= =
48 96
= =
6020 6020
96 15000
or sort -n

sort + uniq = ♥
Pipe something to
'sort | uniq', and you'll
get a deduplicated list
of lines! **sort -u** does the
same thing.
b
b
a
a
| sort -u => b
a
a
or sort | uniq

sort -h: human sort
'sort-n' order 'sort-h' order

15 6 45 K
= 30 M 30 M
= 45 K
= 200 G 15 6
= 200 G 200 G
=

useful example:
du -sh * | sort -h

uniq -c

counts each line it saw.

Recipe: get the top 10 most common lines in a file:
\$ sort foo.txt
| uniq -c
| sort -n
| tail -n 10



Sed

Some more sed incantations...

sed -n 12 p

print 12th line

sed -n suppresses output so only what you print with 'p' gets printed

sed 6
double space a file
(good for long error lines)

sed '/cat/a dog'
append 'dog' after lines

sed -n 5,30 p
delete lines matching /cat/

print lines 5-30

sed 5d
delete 5th line

sed /cat/d
can be a regular expression

sed s+cat/+dog/+
use '+' as a regex delimiter
way easier than escaping '/' is like
s/cat\\//dog\\/!

sed -n s/cat/dog/p
only print changed lines
insert "panda" on line 17

bash tricks

10

★ ctrl + r ★ search your history! I use this constantly to rerun commands	★ magical braces★ \$ convert file.{jpg,png} expands to \$ convert file.jpg file.png {1..5} expands to 1 2 3 4 5 (for i in {1..100} ...)	loops for i in *.png do convert \$i \$i.jpg done	for loops: easy & useful!
\$ convert file.jpg file.png {1..5} expands to 1 2 3 4 5 (for i in {1..100} ...)	loops \$() gives the output of a command \$ touch file-\$(date -I) ↗ create a file named file-2018-05-25	more keyboard shortcuts ctrl a beginning of line ctrl e end of line ctrl l clear the screen + lots more emacs shortcuts too!	more keyboard shortcuts ctrl a beginning of line ctrl e end of line ctrl l clear the screen + lots more emacs shortcuts too!

top	load average 3 numbers that roughly reflect demand for your CPUs on the system in the last 1, 5, and 15 minutes if it's higher than the # of CPUs you have, that's often bad	top a live-updating summary of the top users of your system's resources	% CPU 350%? what?
15	memory 4 numbers: total / free / used / cached One perhaps unexpected thing: total is <u>not</u> free + used! total = free + used + cached filesystem cache	memory 4 numbers: total / free / used / cached One perhaps unexpected thing: total is <u>not</u> free + used! total = free + used + cached filesystem cache	htop a prettier & more interactive version of top

PS

ps

ps shows which processes are running

I usually run ps like this:

```
$ ps aux
u means include username column
a+* together show all process
```

(ps -ef works too)

process state

Here's what the letters in ps's STATE column mean:

- R: running
- S/D: asleep
- Z: zombie
- t: multithreaded
- +: in the foreground

w is for wide: ps auxww will show all the command line args for each process

e

is for environment: ps auxe will show the environment vars!

wchan you can choose which columns to show with ps (-eo ...). One cool column is 'wchan' which tells you the name of the kernel function if the process is sleeping. try it:

```
ps -eo user,pid,wchan,cmd
```

f is for "forest": ps auxf will show you an ASCII art process tree!

pstree can display a process tree too

ps has 3 different sets of command line arguments ↪

1. UNIX (1 dash)
2. BSD (no dash)
3. GNU (2 dashes)

You can write monstrously like:

```
$ ps f -f
↑ full format (BSD)
forests (UNIX)
```

more bash tricks

11

cd -

changes to the directory you were last in

pushd & popd let you keep a stack

<()

process substitution treat process output like a file (no more temp files!)

eg:

```
$ diff <(ls) <(ls -a)
```

ctrl+z

suspends (SIGTSTP) the running program

fg

brings background/suspended program to the foreground

bg

starts suspended program & backgrounds it (use after ctrl+z)

shellcheck

shell script linter! helps spot common mistakes.

type

tells you if something is a builtin, program, or alias

try running type on [time] [ping] [pushd] (they're all different types)

disk usage

12

du

tells you how much disk space files / directories take up
-S summary: total size of all files in a directory
-h human readable sizes

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/sda3	18G	6	2.5G	86%	/
udev	483M	4.0K	483M	1%	/dev
tmpfs	99M	1.4M	97M	2%	/run
/dev/sda4	167G	157G	9.9G	95%	/home

df -i

instead of % disk free, report how many **inodes** are used/ free on each partition



ncdu

see what's using disk space navigate with arrow keys

17.5 GiB [####]	/music
3.2 GiB [##]	/pictures
5.7 MiB [/text
2.0 MiB [/file.pdf

iostat

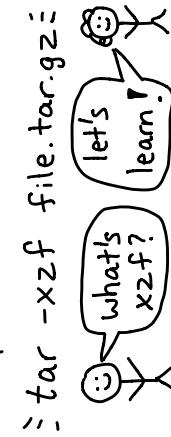
get statistics about disk reads / writes interval to report at

```
# iostat 5
Device: kB_read/s kB_wrtn/s
sda      2190.21   652.87
sdb       6.00     0.00
```

tar

13

The .tar file format combines many files into one file.
.tar files aren't compressed by themselves. Usually you gzip them: .tar.gz or .tgz!



Usually, when you use the 'tar' command, you'll run some incantation To unpack a tar.gz, use:

```
-tar -xzf file.tar.gz
```

-t is for list lists the contents of a tar archive

-f is for file which tar file to create or unpack & more! see the man page ↴

tar can compress / decompress

- Z gzip format (.gz)
- j bzip2 format (.bz2)
- J xz format (.xz)

putting it together

- list contents of a .tar.bz2: \$ tar -tzvf file.tar.bz2 verbose
- create a .tar.g2: \$ tar -czf file.tar.gz dir/ files to go in the archive