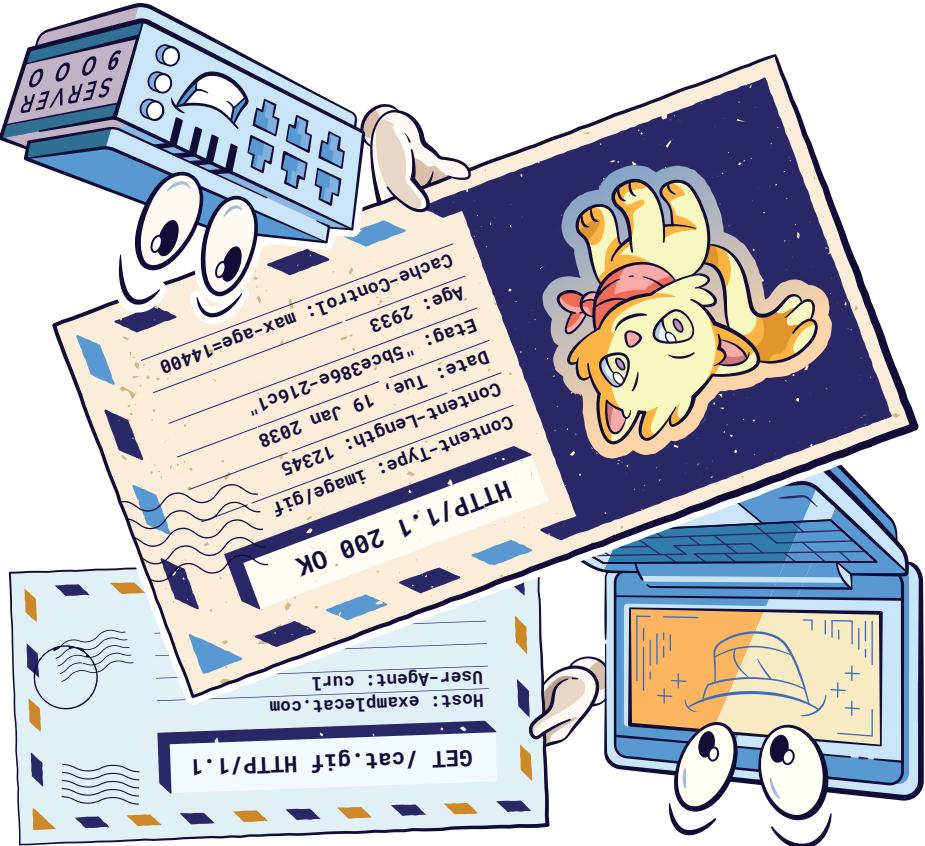


by Julia Evans



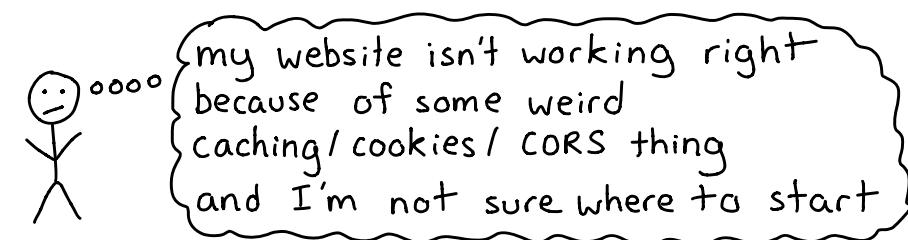
HTTP
Learn your browser's language

Like this?
more zines at
wizardzines.com

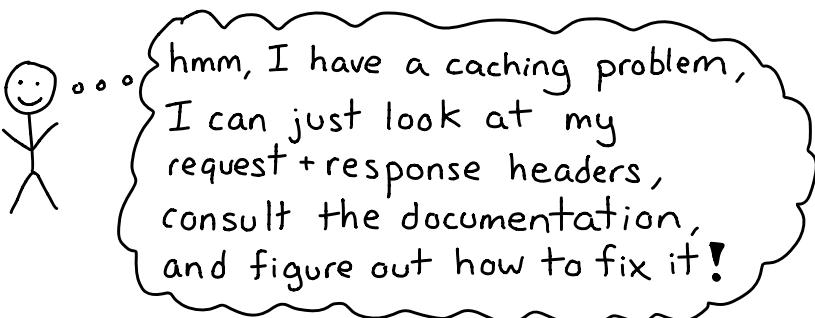
about this zine

Your browser uses HTTP every time it visits a website. Like a lot of the tech that runs the internet, understanding HTTP isn't that hard!

This zine's goal is to take you from:



to



credits

Cover art: Vladimir Kašiković
Editing: Dolly Lanuza, Kamal Marhubi
special thanks to Marco Rogers for suggesting the idea of a HTTP zine

how to learn more

♥ Mozilla Developer Network

<https://developer.mozilla.org>

MDN is a fantastic wiki maintained by Mozilla. It has tutorials and reference documentation for HTML, CSS, HTTP, Javascript. It's the best place to start for reference documentation.

♥ OWASP

<https://cheatsheetseries.owasp.org>

OWASP is an organization that publishes security best practices. If you have a question about web security, they've probably published a cheat sheet or guide to help you.

♥ httpstatuses.com

Nice little site that explains all the HTTP status codes.

♥ RFCs

<https://tools.ietf.org/html/rfcXXXX>

put RFC number here

RFCs are numbered documents (like "RFC 2631"). Every Internet protocol (like TLS or HTTP) has an RFC. These are where you go to find the Official Final Answers to technical questions you have about any internet standard. The HTTP standard is mostly documented in 6 RFCs numbered 7230 to 7235.



is the Host header
actually required?



Yes, section 5.4
of RFC 7230
says so!

the final answer
Don't be scared of using an RFC if you want to know for sure!

What's HTTP?	4
How URLs work	5
What's a header?	6
Request methods	7
Anatomy of an HTTP request	8-9
Request methods (GET! POST!)	8-9
Using HTTP APIs	10
Request headers	11
Responses:	12
Anatomy of an HTTP response	13
Status codes (200! 404!)	14
How cookies work	15
Content delivery networks & caching	16-17
HTTP/2	18
Redirections	19
HTTP/2	20-21
Same origin policy & CORS	22-24
Security headers	25
Exercises & how to learn more	26-27

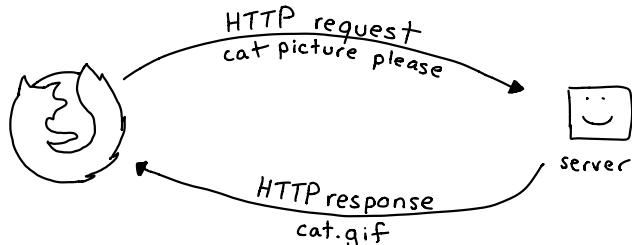
Table of Contents

Making HTTP requests with curl to real internet websites and trying different headers is my favourite way to play around with HTTP & learn.	
curl -i shows the response headers by sending a HEAD request	-i shows the response headers
-H adds a request header	-H adds a request header
Try the Range header:	curl -i https://examplecat.com/cat.txt -H "Range: bytes=8-17"
Request (and print out!) a compressed response:	curl -i https://examplecat.com/cat.txt -H "Accept-Encoding: gzip" -- output -
Get a webpage in Spanish:	curl -i https://twitter.com -H "Accept-Language: es-ES"
Get redirected to another URL:	(hint: look at the Location header!)
curl -i http://examplecat.com	curl -i https://github.com/itsinahader
Guess what content delivery network GitHub is using:	(hint: it's in a header starting with x-)
curl -I https://github.com/bananas	curl -I https://github.com/last-modified
Find out when example.com was last updated	curl -I example.com
Get a 404 not found	curl -I example.com
curl -i example.cat	curl -i example.cat
Get a 404 not found	curl -I example.cat
curl -i example.cat.com	curl -I example.cat.com
curl -i example.cat.com/bananas	curl -I example.cat.com/bananas

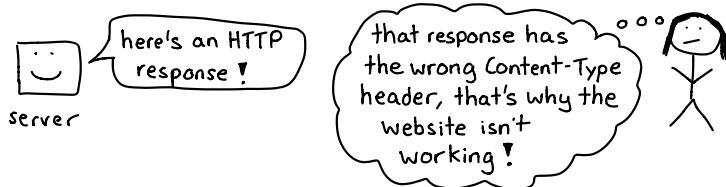
HTTP exercises

what's HTTP?

HTTP is the protocol (**Hypertext Transfer Protocol**) that's used when you visit any website in your browser.



The exciting thing about HTTP is that even though it's used for literally every website, HTTP requests and responses are easy to look at and understand:



Example of what an HTTP request and response might look like:

request	response
request line headers { Host: examplecat.com User-Agent: curl Accept: */*	status { HTTP/1.1 200 OK headers { Cache-Control: max-age=604800 Content-Type: text/html Etag: "1541025663+ident" Server: ECS (nyb/1D0B) Vary: Accept-Encoding X-Cache: HIT Content-Length: 1270
body {	<!doctype html> <title>Example Cat</title> ...

All that text is a lot to understand, so let's get started learning what all of it means!

security headers

These are headers your server can set. They ask the browser to protect your users' data against attackers in different ways:

Content-Security-Policy often called CSP

Only allow CSS / Javascript from certain domains you choose to run on your website. Helps protect against cross-site-scripting (aka XSS) attacks.

Referrer-Policy

Control how much information is sent to other sites in the Referer header. Example: Referrer-Policy: no-referrer.
spelling is inconsistent with Referer header !!

Strict-Transport-Security often called HSTS

Require HTTPS. If you set this the client (browser) will never request a plain HTTP version of your site again. Be careful! You can't take it back!

Expect-CT

Certificate Transparency (CT) is a system that can help find malicious SSL certificates issued for your site. This header gives the browser a URL to use to report bad certificates to you.

X-XSS-Protection

Another way to protect against XSS attacks. Not supported by all browsers, Content-Security-Policy is more powerful.

How URLs work

https://example.com:443/cats?color=light%20gray#banana
scheme domain port path query string fragment id

Protocol to use for the requests. Encrypted (https), or something else entirely (ftp). Insecure (http), or something else entirely (ftp).

Protocol to use for the request. Encrypted (https), **scheme https://** **insecure (http), or something else entirely (ftp).**

Defaults to 80 for HTTP and 443 for HTTPS.

Query parameters are usually used to ask for a different version of a page ("I want a light

hair=short&COLOR=black&name=mir%20darci
name = value separated by &

This isn't sent to the server at all. It's used either to jump to an HTML tag (``) or by `javaScript` on the page.

% + hex representation of ASCII value.

characters like spaces, @, etc. So to put URL you need to percent encode them as

URLs aren't allowed to have certain special

```
name = value separated by &  
hair=short&color=black&name=mrsودادری
```

a different version of a page ("I want a light gray cat"). Example:

query parameters are combined in the request, like: GET /cats?color=light%20gray HTTP/1.1

Path to ask the server for. The path and the

Defaults to 80 for HTTP and 443 for HTTPS.

the Host header gets set to this (Host: example.com).

Where to send the request. For `HTTP(s)` requests

Protocol to use for the request. Encrypted (<https://>) or something else entirely (<ftp://>).

`https://example.com`

<https://scheme>

490

/cats

parameters

#banana

This OPTIONS request is called a "preflight" request, and it only happens for some requests, like we described in the diagram on the same-origin policy page. Most GET requests will just be sent by the browser without a preflight request first, but POST requests that send JSON need a preflight.

If you run `api.clothes.com`, you can allow `clothes.com` to make requests to it using the `Access-Control-Allow-Origin` header.

Here's what happens:

```

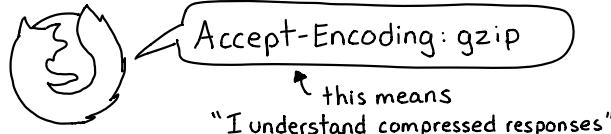
graph LR
    Client((Client)) -- "POST /buy-thing" --> Proxy["Proxy Script"]
    Proxy -- "OPTIONS /buy-thing" --> Server["Host: api.clothes.com"]
    Server -- "hey, what requests are allowed?" --> Response["Response: allowed requests"]
    Response --> Proxy
    Proxy -- "allowed requests" --> Client
  
```

The diagram illustrates a sequence of API requests:

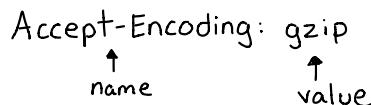
- Client Request:** `Access-Control-Allow-Origin: clothes.com`
- Server Response:** `cool, the request is allowed!`
- Client Request:** `POST /buy-thing`
- Server Response:** `Host: api.clothes.com`
- Client Request:** `Referer: api.clothes.com/checkout`
- Server Response:** `200 OK`
- Client Request:** `{"thing_bought": true}`

what's a header?

Every HTTP request and response has headers. Headers are a way for the browser or server to send extra information!



Headers have a name and a value.



Header names aren't case sensitive:

totally valid → accept-encoding : gzip

There are a few different kinds of headers:

Describe the body:

Content-Type: image/png Content-Encoding: gzip
Content-Length: 12345 Content-Language: es-ES

Ask for a specific kind of response:

Accept: image/png
Range: bytes=1-10

Accept-Encoding: gzip
Accept-Language: es-ES

Every Accept-header has a corresponding Content- header

Manage caches:

ETag: "abc123"
If-None-Match: "abc123"
Vary: Accept-Encoding

If-Modified-Since: 3 Aug 2019 13:00:00 GMT
Last-Modified: 3 Feb 2018 11:00:00 GMT
Expires: 27 Sep 2019 13:07:49 GMT
Cache-Control: public, max-age=300

Say where the request comes from:

User-Agent: curl

Referer: https://examplecat.com

Cookies:

Set-Cookie: name=julia; HttpOnly (server → client)
Cookie: name=julia (client → server)

6

and more!

why the same origin policy matters

Browsers work hard to make sure that evil.com can't make requests to other-website.com. But evil.com can request other-website.com from its own server, what's the big deal?

2 reasons it's important to restrict Javascript on websites from making arbitrary requests from your browser:

Reason 1: cookies

Browsers often send your cookies with HTTP requests. You don't want evil.com to be able to make requests using your login cookies. They'd be logged in as you!



Reason 2: network access

You might be on a private network (for example your company's corporate network) that evil.com doesn't have access to, but your computer does.

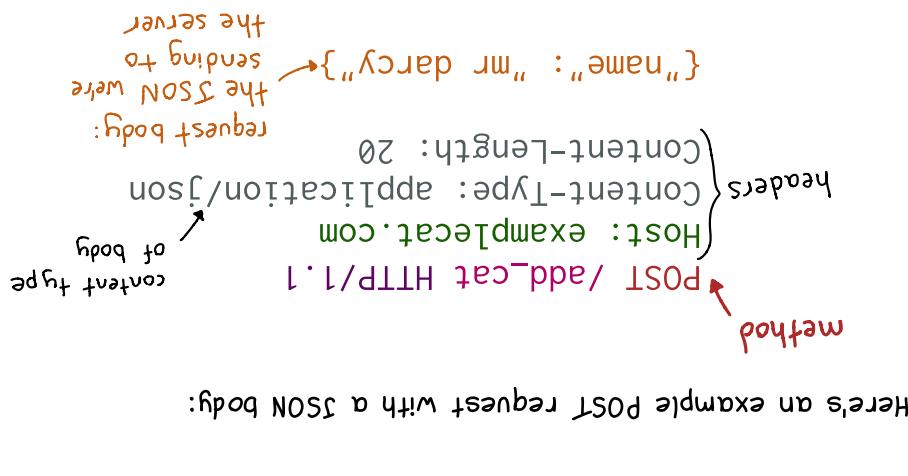


23

* Anatomy of an HTTP request *

HTTP requests always have:

- a domain (like example.com)
- a resource (like /cat.jpg)
- a method (GET, POST, or something else)
- headers (extra information for the server)
- There's an optional request body. GET requests usually don't have a body, and POST requests usually do.
- This is an HTTP 1.1 request for example.com/cat.jpg
- It's a GET request, which is what happens when you type a URL in your browser. It doesn't have a body.
- usually GET or POST
- method
- resource being requested
- HTTP version
- domain being requested
- GET /cat.jpg HTTP/1.1
- Host: example.com
- User-Agent: Mozilla...
- Cookies:
- Headers



The same origin policy is one way browsers protect you from malicious JavaScript code. Here's basically how it works:

An origin is the protocol + domain including subdomains + port

example: `https://babby.example.com:443`

Same origin is `http://evil.com`

Please make a request to this URL:

`let me check my following character..`

match exactly?

does the origin

is this request type allowed

from a different origin?

CORS time

notice the default

is no, not yes!

allowed

yes!

few other things are OK
``, and a
`<css>`

is this request type allowed

from a different origin?

yes!

no!

"Is this a simple request?"

for example a GET request with no body and no extra headers

make the request!

yes!

no!

do the response's CORS headers allow this?

do an OPTIONS preflight request

headers allow this?

yes!

no!

DENIED

ALLOWED

request methods

Every HTTP request has a method. It's the first thing in the first line:

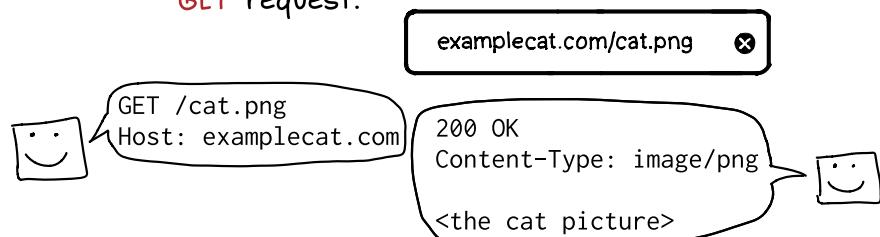
this means it's a GET request

GET /cat.png HTTP/1.1

There are 9 methods in the HTTP standard. 80% of the time you'll only use 2 (GET and POST).

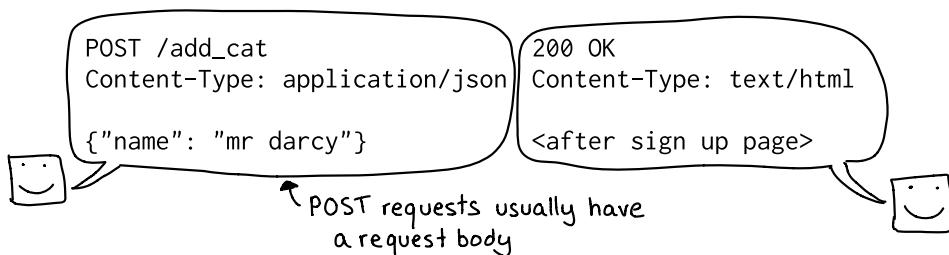
GET

When you type an URL into your browser, that's a GET request.



POST

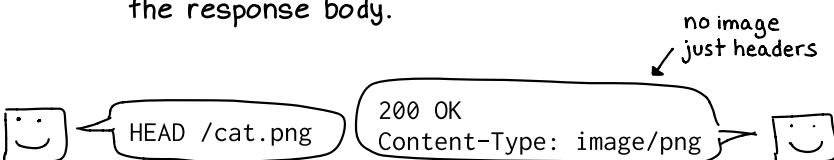
When you hit submit on a form, that's (usually) a POST request.



The big difference between **GET** and **POST** is that **GET**s are never supposed to change anything on the server.

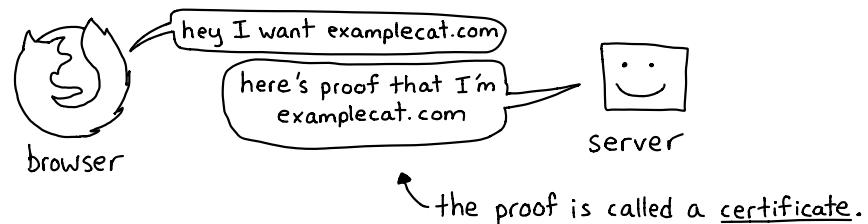
HEAD

Returns the same result as GET, but without the response body.



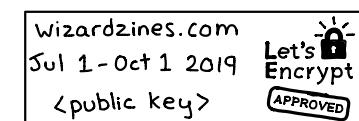
certificates

To establish an HTTPS connection to examplecat.com, the client needs proof that the server actually is examplecat.com.



A TLS certificate has:

- a set of domains it's valid for (eg examplecat.com)
- a start and end date (example: july 1 2019 to oct 1 2019)
- a secret private key which only the server has this is the only secret part, the rest is public
- a public key to use when encrypting
- a cryptographic signature from someone trusted



The trusted entity that signs the certificate is called a ★ Certificate Authority ★ (CA) and they're responsible for only signing certificates for a domain for that domain's owner.



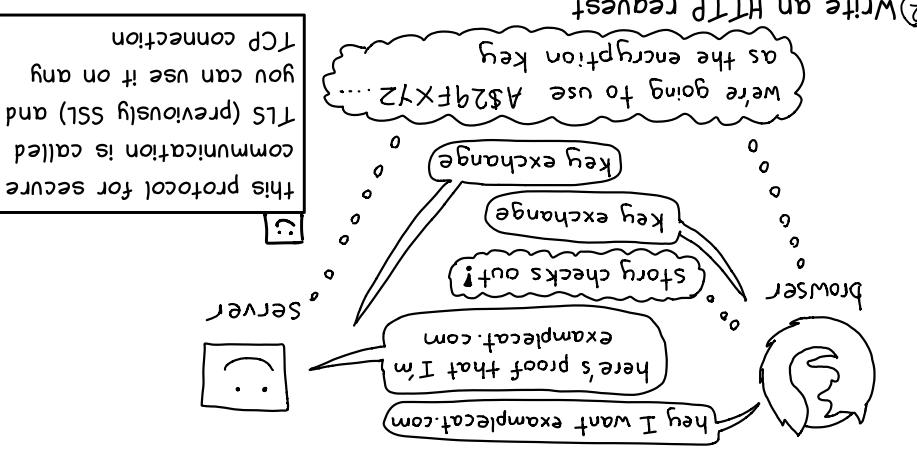
When your browser connects to examplecat.com, it validates the certificates using a list of trusted CAs installed on your computer. These CAs are called "root certificate authorities".



HTTPS: HTTP + Secure ⚡

Here's what your browser does when it asks for `https://examplecat.com/cat.png`:

① Negotiate an encryption key (AES symmetric key) to use for this connection to `examplecat.com`. The browser and server will use the same key to encrypt/decrypt content.



Simplified version of how picking the encryption key works:

server will use the same key to encrypt/decrypt content.

use for this connection to `examplecat.com`. The browser and



③ Encrypt the HTTP request with AES & send it to `examplecat.com`

User-Agent: Mozilla/5.0
Host: examplecat.com
GET /cat.png HTTP/1.1

② Write an HTTP request

as the encryption key
we're going to use A\$24FX7Z
TCP connection
this protocol for secure
communication is called
TLS (previously SSL) and
you can use it on any

CONNECT

If you set the `HTTPS_PROXY` environment variable to a proxy server, many HTTP libraries will use this protocol to proxy your requests.

TRACE

I've never seen a server that supports this, you probably don't need to know about it.

PATCH

Used in some APIs for partial updates to a resource ("just change this file").

PUT

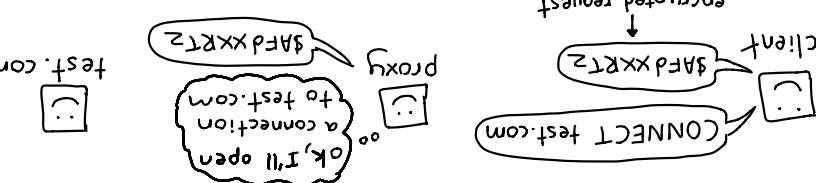
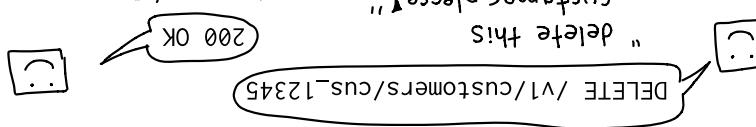
Used in some APIs (like the S3 API) to create or update resources. `PUT /cat/1234` lets you `GET /cat/1234` later.

OPTIONS

It also tells you which methods are available. The CORS page has more about that. `OPTIONS` is mostly used for CORS requests.

DELETE

Used in many APIs (like the Stripe API) to delete resources.



request headers

These are the most important request headers:

Host

The domain.
The only required header.

Host: examplecat.com User-Agent: curl 7.0.2 Referer: https://examplecat.com

User-Agent

Name + version of your browser and OS

User-Agent: curl 7.0.2 Referer: https://examplecat.com

Referer

website that linked or included the resource

Referer: https://examplecat.com
↑ yes, it's misspelled!

Authorization

eg a password or API token
base64 encoded user:password

Authorization: Basic YXZ

Cookie

Send cookies the server sent earlier
keeps you logged in.

Cookie: user=b0rk

Range

lets you continue downloads ("get bytes 100-200")

Range: bytes=100-200

Cache-Control

"max-age = 60"
means cached responses must be less than 60 seconds old

If-Modified-Since: Wed, 21 Oct...

Accept-Encoding

set this to "gzip" and you'll probably get a compressed response

Accept-Encoding: gzip

Accept-Language

set this to "fr-CA" and you might get a response in French

Accept-Language: fr-CA

Content-Type

MIME type of request body, e.g. "application/json"

Content-Encoding

will be "gzip" if the request body is gzipped

Connection

"close" or "keep-alive". Whether to keep the TCP connection open.

HTTP/2

HTTP/2 is a new version of HTTP.

Here's what you need to know:

★ A lot isn't changing

All the methods, status codes, request/response bodies, and headers mean exactly the same thing in HTTP/2.

before (HTTP/1.1)

method: GET
path: /cat.gif
headers:
- Host: examplecat.com
- User-Agent: curl

after (HTTP/2)

method: GET
path: /cat.gif
headers:
- User-Agent: curl
authoritative: examplecat.com

★ HTTP/2 is faster

Even though the data sent is the same, the way HTTP/2 sends it is different. The main differences are:

- It's a binary format (it's harder to tcpdump traffic and debug)
- Headers are compressed
- Multiple requests can be sent on the same connection at a time

before (HTTP/1.1)

→ request 1
→ request 2
→ response 1 ←
→ response 2 ←

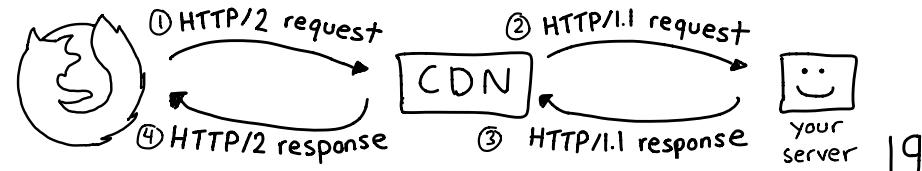
after (HTTP/2)

→ request 1
→ request 2
out of order { response 2 ←
order is ok } response 1 ← } one TCP connection

All these changes together mean that HTTP/2 requests often take less time than the same HTTP/1.1 requests.

★ Sometimes you can switch to it easily

A lot of software (CDNs, nginx) let clients connect with HTTP/2 even if your server still only supports HTTP/1.1.



Using HTTP APIs

lots of services (Twitter! Willow! Google!) let you use them by sending them HTTP requests. If an HTTP API doesn't come with a client library, don't be scared! You can just make the HTTP requests yourself. Here's what you need to remember:

Often you'll be sending a POST request with a body, and that means you need a Content-Type header that matches the body.

- * application/json → JSON;
- * application/x-www-form-urlencoded → same as what a HTML form does

a common error is to try to send POST data as one content type (like JSON) when it's actually another like application/x-www-form-urlencoded

- Identify yourself
- Most HTTP APIs require a secret API key so they know where's how that looks for the Twilio API:
`curl https://api.twilio.com/2010-04-01/Accounts/ACCOUNT_ID/Me/H_Content-Type: application/json`
- u ACCOUNT_ID: AUTH_TOKEN
-d {
 - u sends the user info
 - "from": "+15141234567",
"to": "+15141234567",
"body": "a text message"
- in the Authorization header
- this sends a POST request

Redirections

Sometimes you type a URL into your browser:

but end up at a slightly different URL:
ooh, where did the cat come from?

 examplecat.com/cat.png

A sequence diagram illustrating a client interaction with a server through a proxy. The client sends a `GET /dog.png` request to the proxy, which returns a `301 Moved Permanently` response with a `Location: /cat` header. The client then sends a `GET /cat.png` request to the proxy, which returns a `200 OK` response with a `Host: examplecat.com` header. The proxy also includes a `Server: browser` header in its responses.

```
sequenceDiagram
    participant Client
    participant Proxy
    participant Server

    Client->>Proxy: GET /dog.png
    Note over Client: Location: /cat
    Note over Client: 301 Moved Permanently
    Proxy-->>Client: 200 OK
    Client->>Proxy: GET /cat.png
    Note over Client: Host: examplecat.com
    Note over Client: 200 OK
    Proxy-->>Client: 200 OK
    Note over Client: Server: browser
```

The location header tells the browser what new URL to use. The new URL doesn't have to be on the same domain: `exampleLocat.com/panda` can redirect to `pandas.com`. Setting up redirects is a great thing to do if you move your site to a new domain!

301 Moved Permanently redirects are PERMANENT: after a browser sees one once, it'll always use `exampleLocat.com/cat.png` when someone types `exampleLocat.com/cat.png` forever. You can't take it back and decide to not to redirect. If you're not sure you want to redirect your site for eternity, use 302 Found to redirect instead.

! Warning!

Warning!

anatomy of an HTTP response

HTTP responses have:

- a status code (200 OK! 404 not found!)
- headers
- a body (HTML, an image, JSON, etc)

Here's the HTTP response from examplecat.com/cat.txt:

```

HTTP/1.1 200 OK status
Accept-Ranges: bytes
Cache-Control: public, max-age=0
Content-Length: 33
Content-Type: text/plain; charset=UTF-8
Date: Mon, 09 Sep 2019 01:57:35 GMT
Etag: "ac5affa59f554a1440043537ae973790-ssl"
Strict-Transport-Security: max-age=31536000
Age: 0
Server: Netlify

\   /
) ( ' ) cat! ^
( / )
\(_)_|
    
```

} status code

} headers

} body

There are a few kinds of response headers:

when the resource was sent/modified:

Date: Mon, 09 Sep 2019 01:57:35 GMT
Last-Modified: 3 Feb 2017 13:00:00 GMT

about the response body:

Content-Language: en-US Content-Type: text/plain; charset=UTF-8
Content-Length: 33 Content-Encoding: gzip

caching:

ETag: "ac5affa..." Age: 255
Vary: Accept-Encoding Cache-Control: public, max-age=0
security: (see page 25)

X-Frame-Options: DENY Strict-Transport-Security: max-age=31536000
X-XSS-Protection: 1 Content-Security-Policy: default-src https:

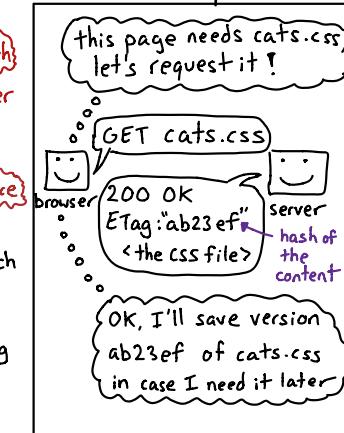
and more:

Connection: keep-alive Accept-Ranges: bytes
Via: nginx
Set-Cookie: cat=darcy; HttpOnly; expires=27-Feb-2020 13:18:57 GMT;

caching headers

These 3 headers let the browser avoid downloading an unchanged file a second time.

initial request



Vary
response header

Sometimes the same URL can have multiple versions (spanish, compressed or not, etc). Caches categorize the versions by request header like this:

Accept-Language	Accept-Encoding	content
en-US	-	hello
es-ES	-	hola
en-US	gzip	f\$xx99aef^.. (compressed gibberish)

The Vary header tells the cache which request headers should be the columns of this table.

Cache-Control

request AND response header

Used by both clients and servers to control caching behaviour. For example:

Cache-Control: max-age=999999999999 from the server asks the CDN or browser to cache the thing for a long time.

Whether Range requests header
is supported for this resource

Accept-Ranges

Whether body is compressed
Content-Encoding: gzip

Content-Encoding

Length of body in bytes
Content-Length: 33

Content-Length

Allow cross-origin requests.
Cross-Origin资源共享头

Access-Control *

Keep-alive
Close or "keep-alive"
TCP connection open

Connection

That response will
vary based on
request headers

Vary

When content was
last modified
(not always accurate)

Last-Modified

response headers

URL to redirect to
Location: /cat.png

Location

Language of body
Content-Language: en-US

Content-Language

MIME type of body
Content-Type: text/plain

Content-Type

Sets a cookie.
Set-Cookie: name=value; HttpOnly

Set-Cookie

The response is stable
and should be re-requested
after this time.

Expires

Various caching
settings

ETag

When response
was sent
Date: Mon, 09 Sep 2019...

Age

How many
seconds response
has been cached

Age: 355

Seconds response

Date: Mon, 09 Sep 2019...

When response

Was sent

I want cat picture?

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

No response

How will I pay
for this?

How much
bandwidth?

Now you owe me

me too!

I want cat picture?

HTTP status codes

Every HTTP response has a ★status code★.



There are 50ish status codes but these are the most common ones in real life:

200 OK

} 2xx s mean
★ success ★

301 Moved Permanently

} 3xx s aren't
errors, just
redirects to
somewhere else

302 Found

temporary redirect

304 Not Modified

the client already has the latest
version, "redirect" to that

400 Bad Request

403 Forbidden

API key/OAuth/something needed

404 Not Found

we all know this one :)

429 Too Many Requests

you're being rate limited

500 Internal Server Error

the server code has an error

503 Service Unavailable

could mean nginx (or whatever proxy)

couldn't connect to the server

504 Gateway Timeout

the server was too slow to respond

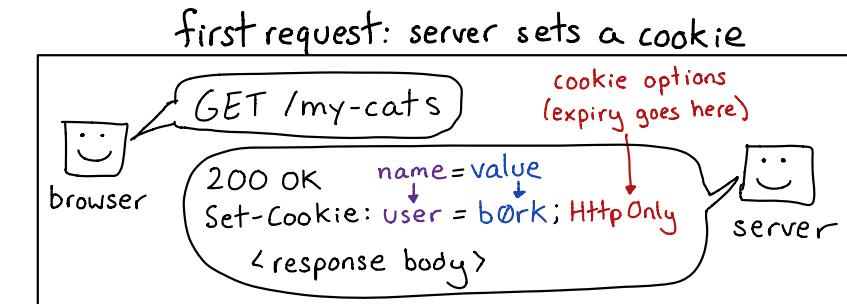
} 4xx errors are
generally the
client's fault:
it made some
kind of invalid
request

} 5xx errors
generally mean
something's wrong
with the server.

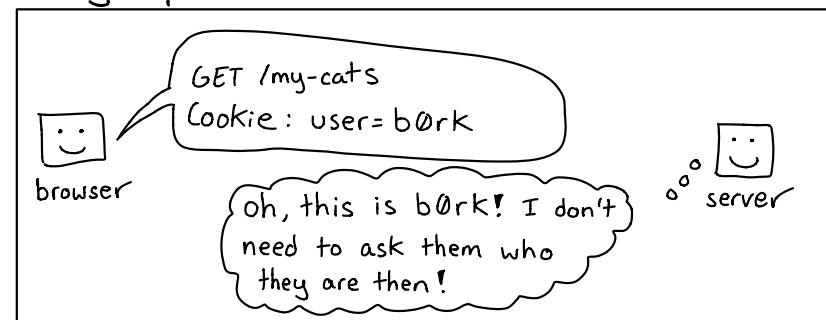
how cookies work

Cookies are a way for a server to store a little bit of information in your browser.

They're set with the Set-Cookie response header, like this:



Every request after: browser sends the cookie back



Cookies are used by many websites to keep you logged in. Instead of user=b0rk they'll set a cookie like sessionid=long-incomprehensible-id. This is important because if they just set a simple cookie like user=b0rk, anyone could pretend to be b0rk by setting that cookie!

Designing a secure login system with cookies is quite difficult — to learn more about it, google "OWASP Session Management Cheat Sheet".